

# Report from the **Crossing the Chasm** Workshop

Report compiled by Mike Rezny,  
Monash University, Australia  
[michael.rezny@monash.edu](mailto:michael.rezny@monash.edu)

Edited and Revised by Bryan Lawrence and  
Reinhard Budich

Report from the ENES Strategy Meeting  
Reading, 26<sup>th</sup> October 2016  
For the IS-ENES2 GA  
16<sup>th</sup> January 2017

Incorporating (hacked) slides  
and ideas from:  
Joerg Behrens,  
Willem Deconinck,  
Rupert Ford,  
Chris Maynard,  
Steve Mullerworth  
Carlos Osuna,  
Andy Porter,  
Kim Serradell  
Sophie Valcke,  
Simon Wilson  
...and the other attendees.

# PROBLEM STATEMENT

## Science Complexity is increasing:

- More components
- More complex interactions between components
- More derived data
- More complex and much larger ensembles
- More observational data for assimilation
- Higher resolution
- More complexity in the grids / meshes
  - Unstructured, extruded meshes
  - Multiple meshes
  - Multi-grid methods
  - High-order, mixed finite element methods

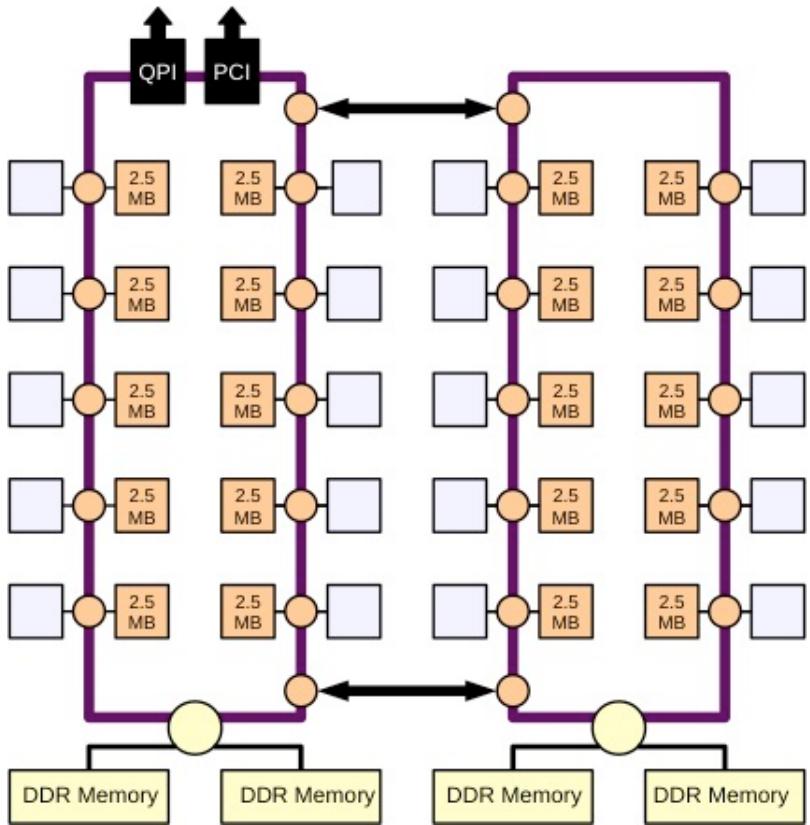
## Boundary Conditions:

- Little or no increase in the performance of each core
- Increased performance is needed to meet research and operational requirements

## Computational Complexity is increasing:

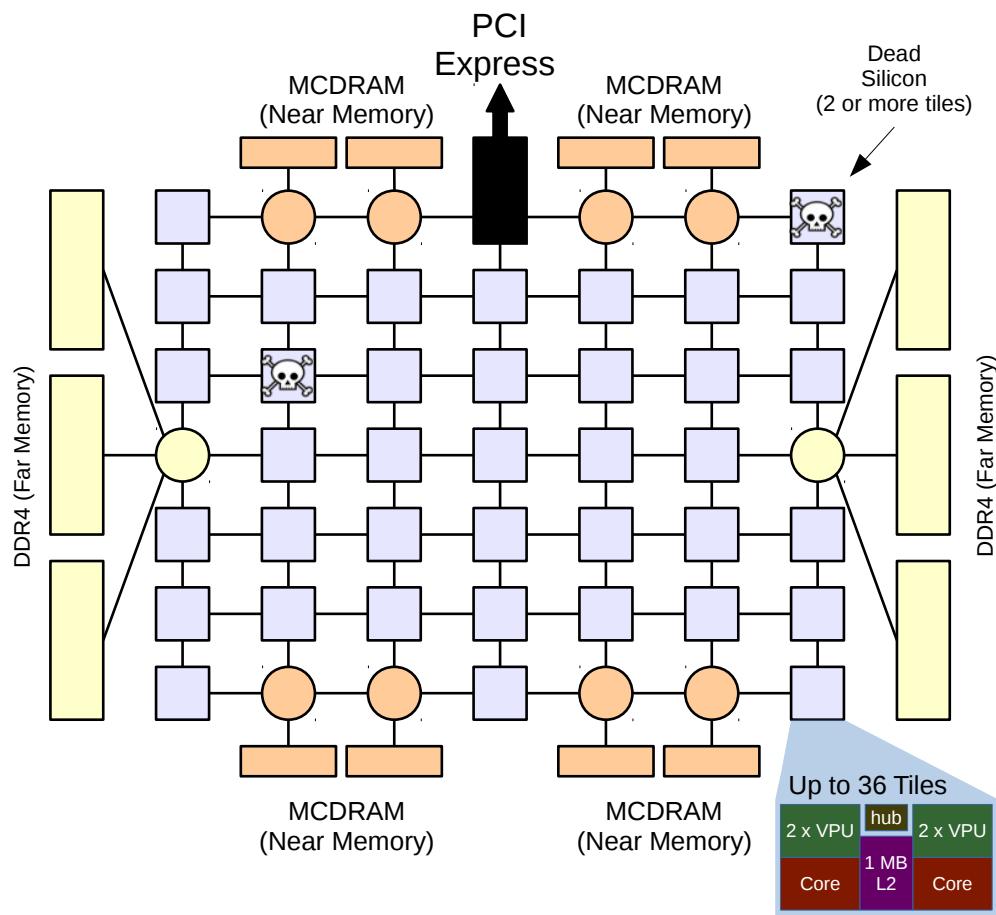
- More cores per socket are provided
- Performance improvement is mainly by higher levels of parallelism
- Heterogeneous hardware architectures with accelerators
- Increasingly complex solutions are being applied to achieve the required performance:
  - Numerics
  - Algorithms
  - Software design
  - Workflows

The difficulty is in applying these complex computational solutions to existing large complex science applications



Cores aren't even the big difference! Memory, cache, and bandwidth differences

# Broadwell and Knights Landing



One Example of Increased Computational Complexity: **Concurrency**

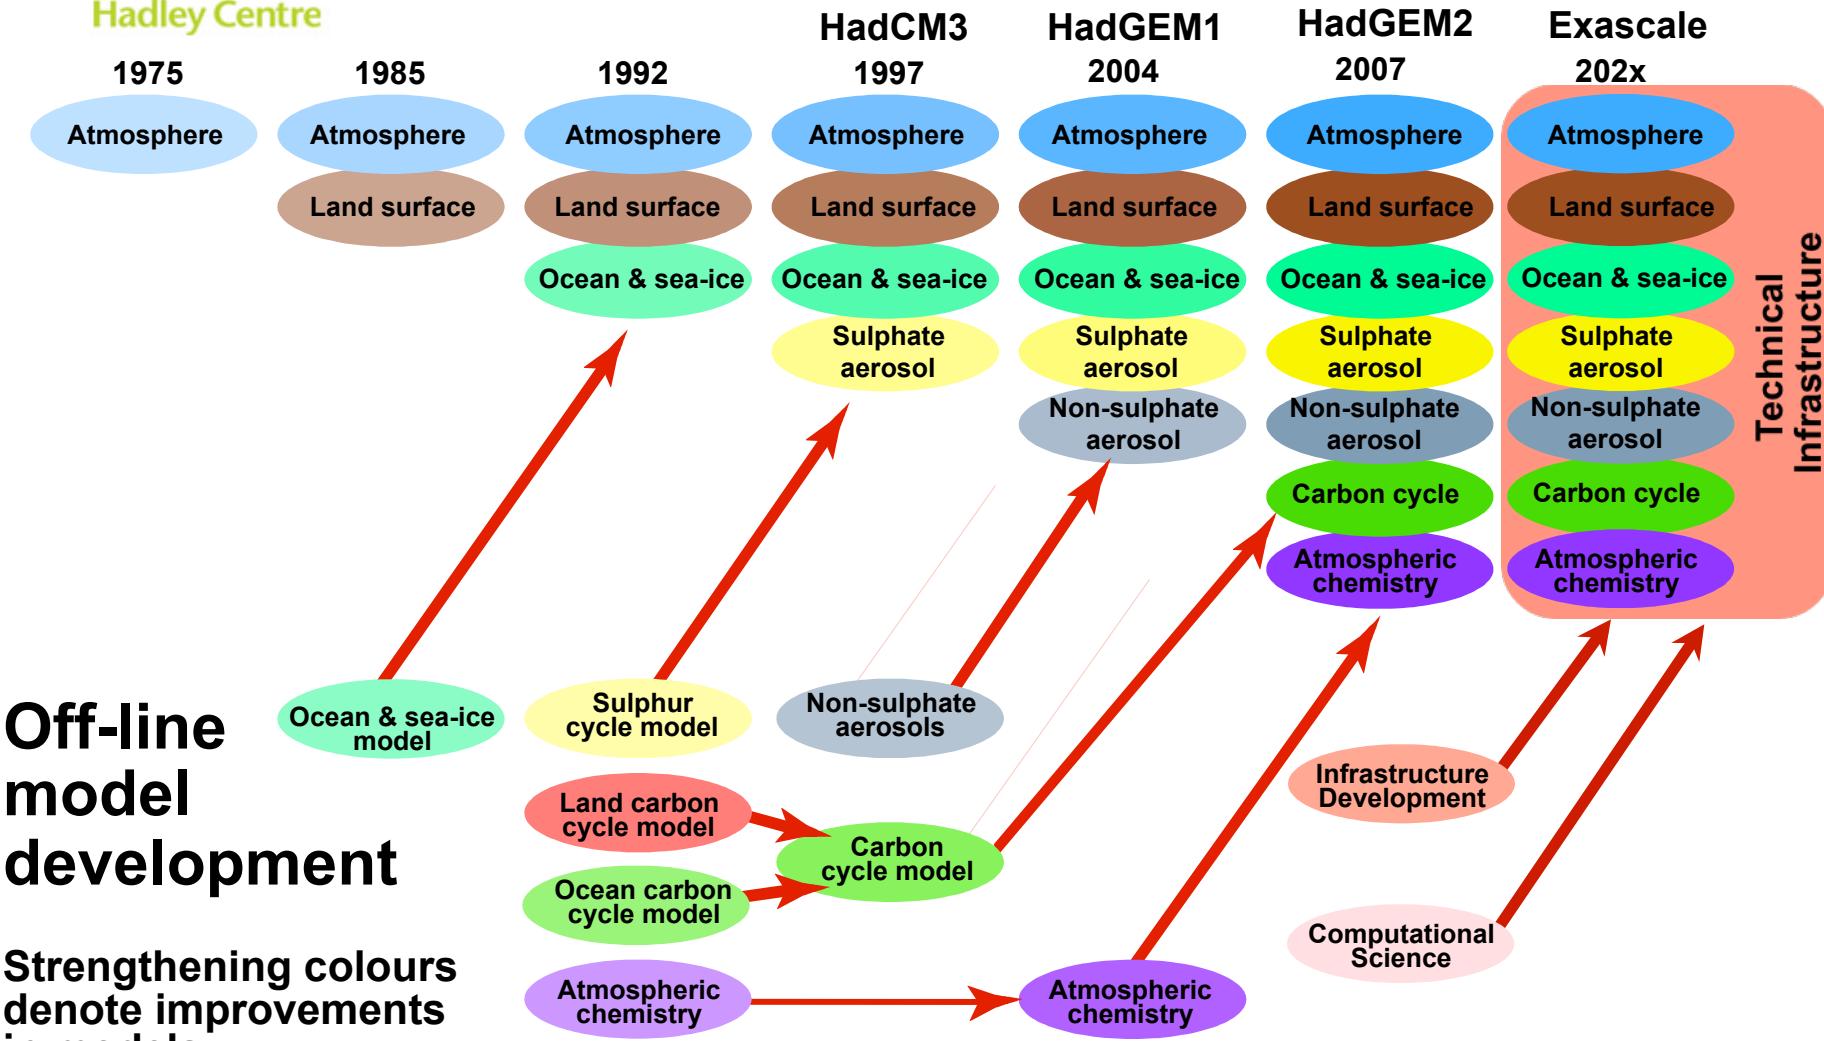
Present and future models may need to incorporate up to **7** types of concurrency to meet performance requirements with anticipated hardware.

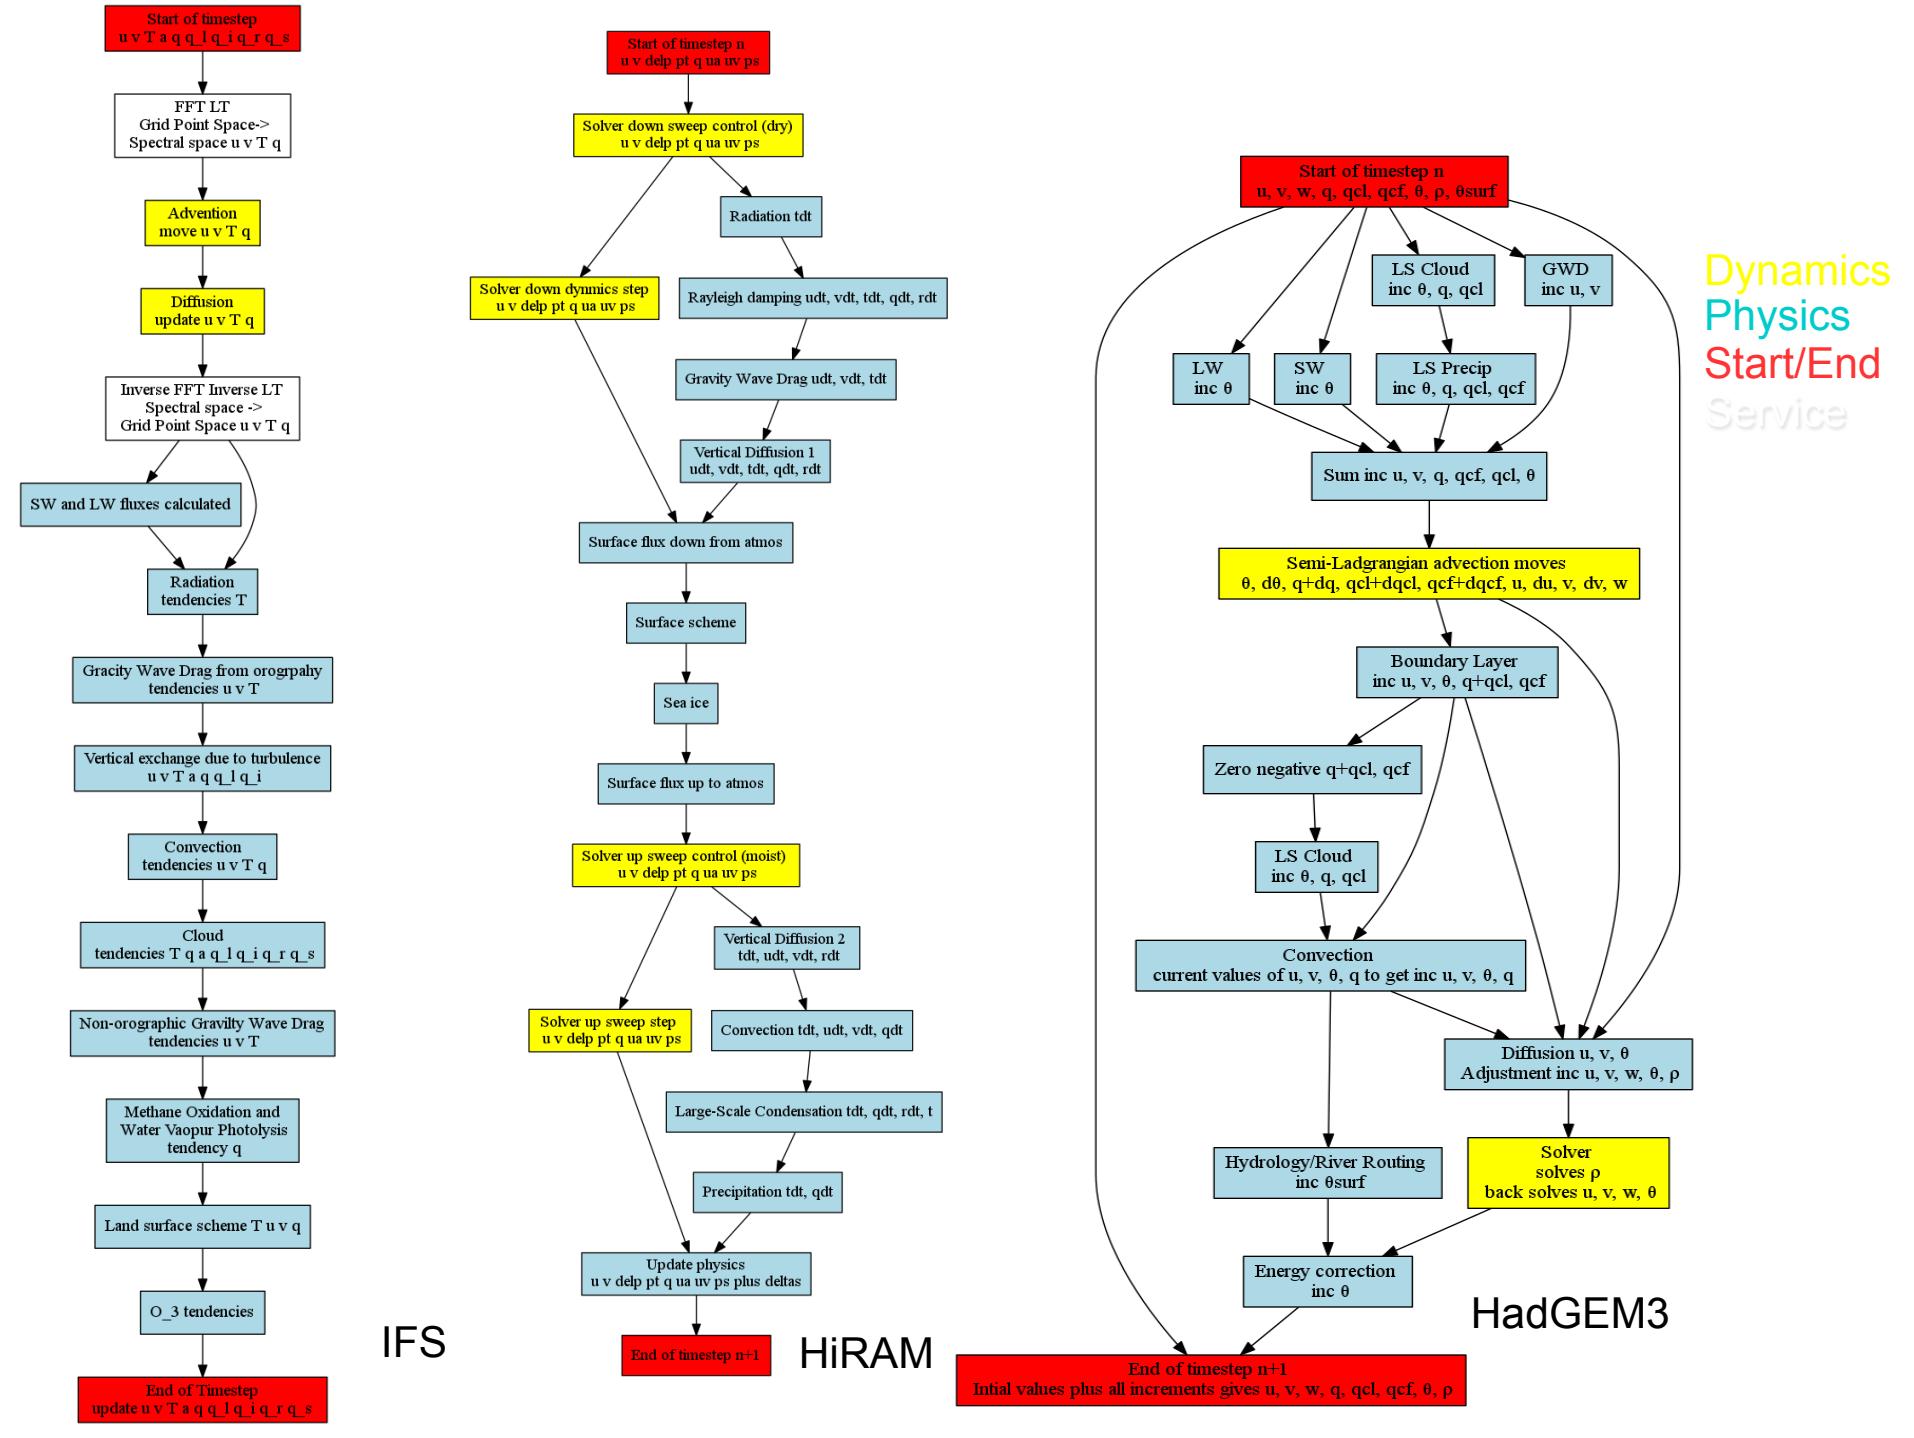
- Vectorisation
- Distributed-memory concurrency across nodes / sockets (MPI)
- Shared-memory concurrency within nodes / sockets (OpenMP)
- External coupling of independent models (Oasis-MCT, Cpl, YAC)
- Asynchronous I/O (XIOS, I/O servers)
- Executing model components concurrently (ESMF, FMS)
- Hardware accelerators (GPU, OpenACC)



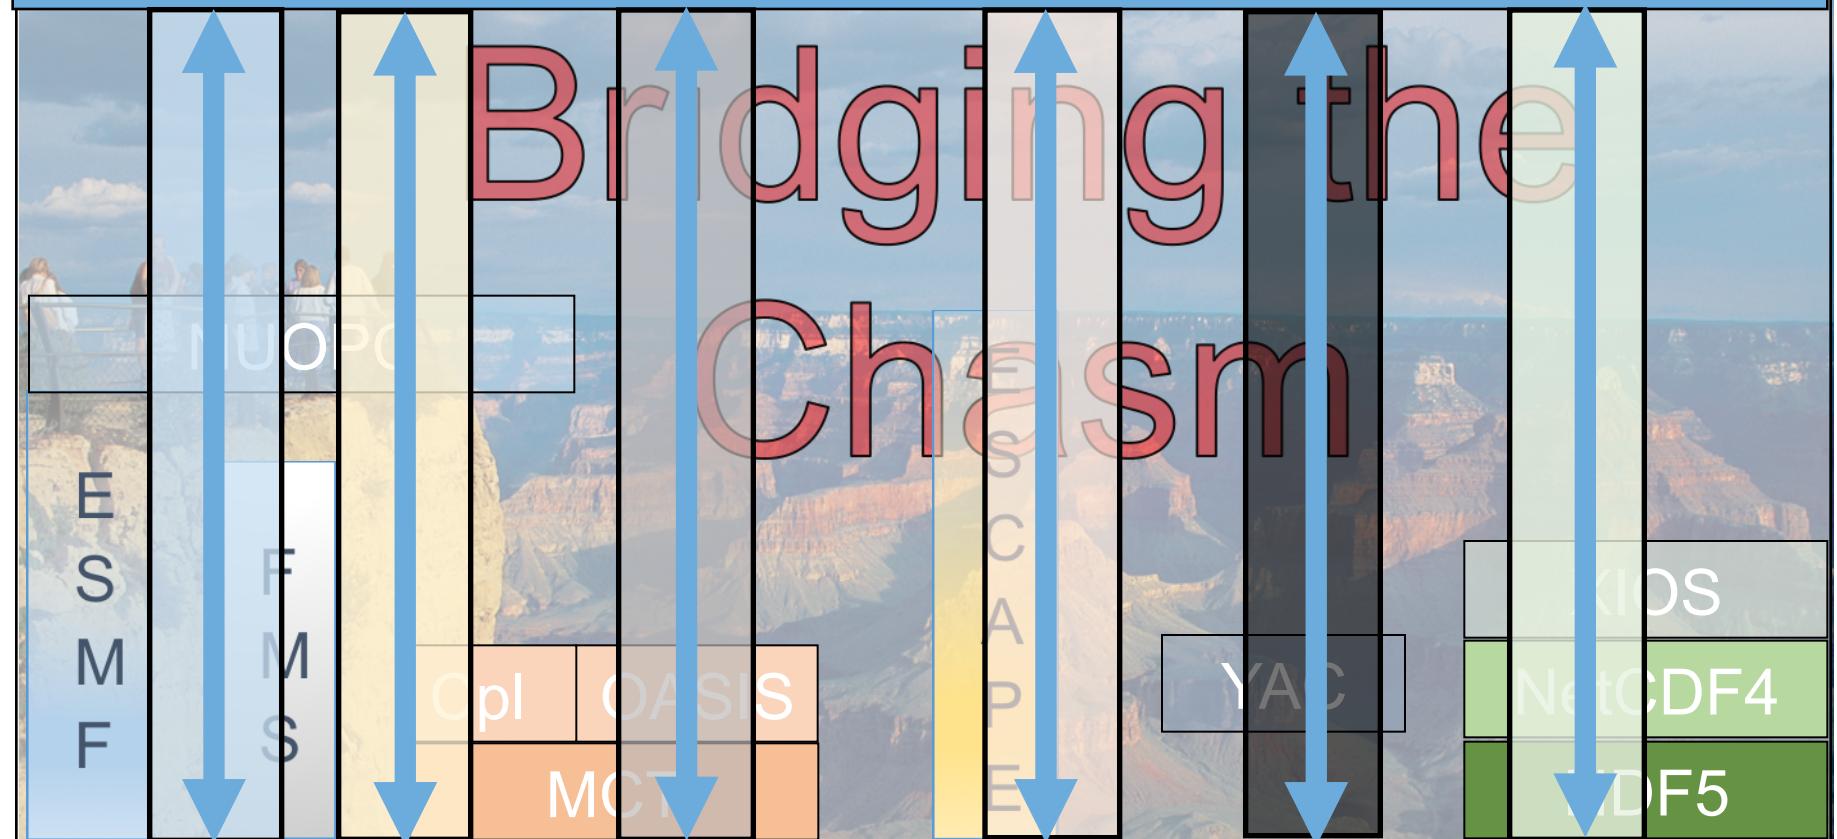
Met Office  
Hadley Centre

# Including Exascale Development Support





# Science Code



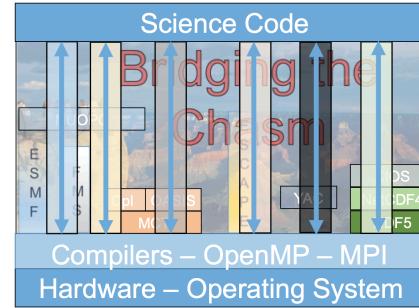
Compilers – OpenMP – MPI  
Hardware – Operating System

Can we, as a community, work collaboratively to:

- Discuss,
- Specify,
- Design,
- Develop,
- Maintain, and
- Document

libraries and tools that will help all of us to  
Progressively Reduce the Chasm in ESM development?

From The top:



# DOMAIN SPECIFIC LANGUAGES



# What is the problem? I

## Codes

- Complex and evolving science
- It is hard to restructure codes
- Software has a long life (20 years +) so outlives hardware architectures

## Software

- Compilers complex and evolving
- Standards evolving (OpenMP, OpenACC, MPI, ...)



# Is there a solution?

## Separation of concerns

- Separate science code from parallelisation and optimisation
- Single science source
- Targeted parallelisation and optimisation for portable performance
- Achievable using domain specific knowledge ... Two important questions:
  - what is a domain?
  - What is the language



Met Office

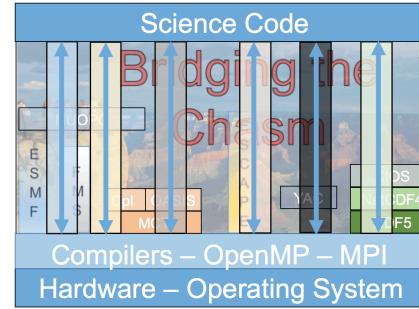
# Experience thus far

1. Separation of concerns and auto-code generation have allowed LFRic to scale to large number of cores (Data Parallelism)
2. Also target shared memory parallelism via OpenMP
3. No science code was altered (or harmed)
4. Vertically contiguous data layout hides indirect memory access cost of unstructured mesh

1. Algorithmic developments (Multi-grid etc) will improve scaling (next target is 221,184 cores 1c)
2. Can switch programming model (OpenACC) for GPU
3. Also looking at Vulcan (ARM) and Knights Landing
4. Look at ILP (kernels) and Task parallelism (driver layer)

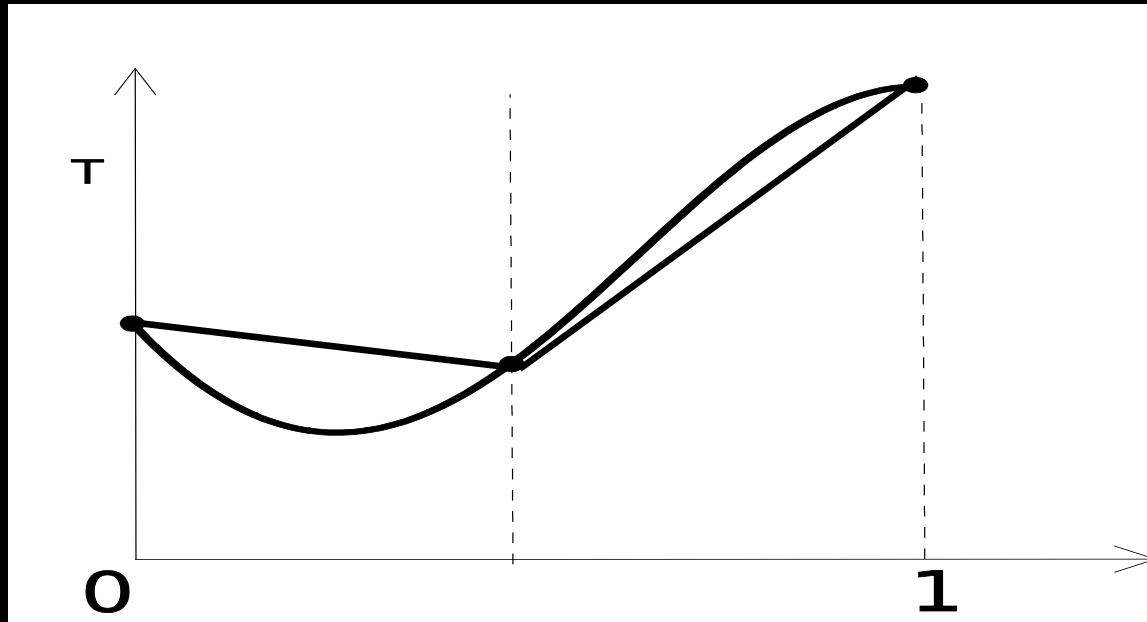
... not dissimilar results with GridTools  
(same idea different implementation)

From The top:



# NEED FOR A DATA MODEL

# What is a data model?

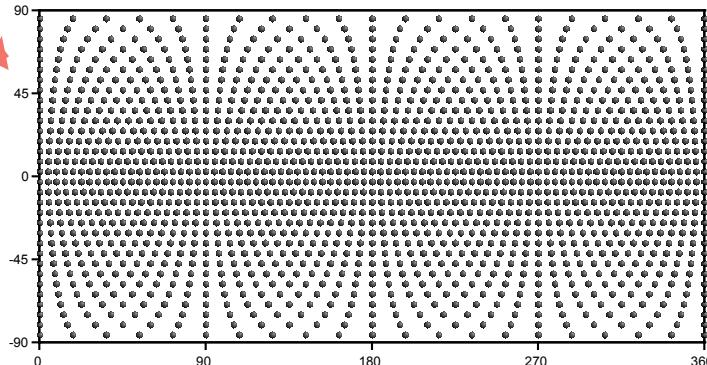
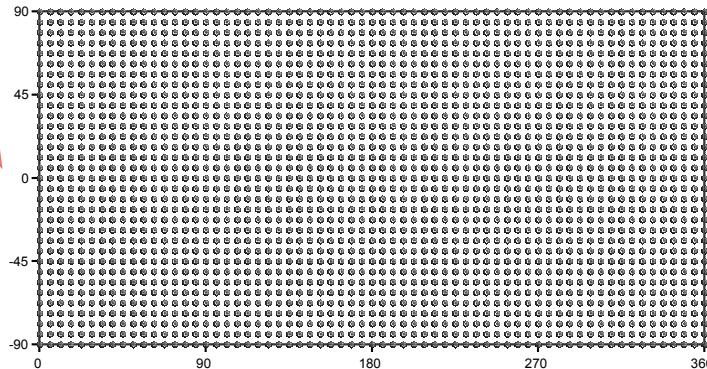
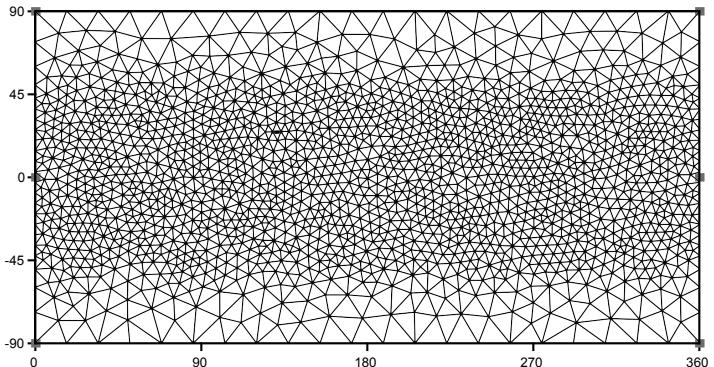


- How do we store the data?
- What happens between data points?
  - Supports science operations and coupling
  - Influences distributed and shared memory deployments

# What is Atlas?

A flexible parallel framework for

- structured grids
- unstructured hybrid meshes



- Parallel mesh generation
- Mathematical operators
- C++ / Fortran

# Planned Atlas developments

- Overlapping communications and computations
- Data storage for many-core architectures
- Parallel interpolation operations
- Spectral/Finite Element basis functions, quadrature
- Serialisation of all data structures to memory (caching)



Met Office  
Hadley Centre

Can we agree on generic data model requirement that meet the need of all ESMs?

Does a one-size fits all approach support both scientific and compute performance goals?

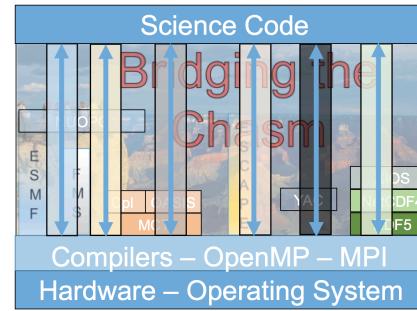
I dunno!  
Je sais pas  
ich weiß nicht



# Summary DSL thoughts

- To cross the chasm we should aim for
  - single source science codes (single source, multiple architectures),
  - higher level specifications and
  - performance portability
- DSEL's are a potential way forward
  - Stella/GridTools and PSyclone/LFRic provide demonstrators
  - Mostly complimentary work
  - Perhaps PSyclone could use Stella?
- Potential for collaboration with optimisations?
  - Convince other groups to try these out
  - Absolute need to involve VENDORS
- Data layout is still a potential issue
  - (data models, un-structured grids)

In the middle?:



# UNDERSTAND THE BUILDING BLOCKS



## ESCAPE key objectives

- Define fundamental **algorithm building blocks** (“*Weather & Climate Dwarfs*”) to co-design, advance, benchmark and efficiently run the next generation of NWP and climate models on energy-efficient, heterogeneous HPC architectures.
- Combine frontier research on **algorithm development** and extreme-scale, high-performance computing applications with **novel hardware technology**, to create a flexible and sustainable weather and climate prediction system.
- Foster the **future design of Earth-system models** and commercialisation of weather-dependent innovative products and services in Europe through enabling open-source technology.
- Pairing **world-leading NWP** with **innovative HPC solutions**.



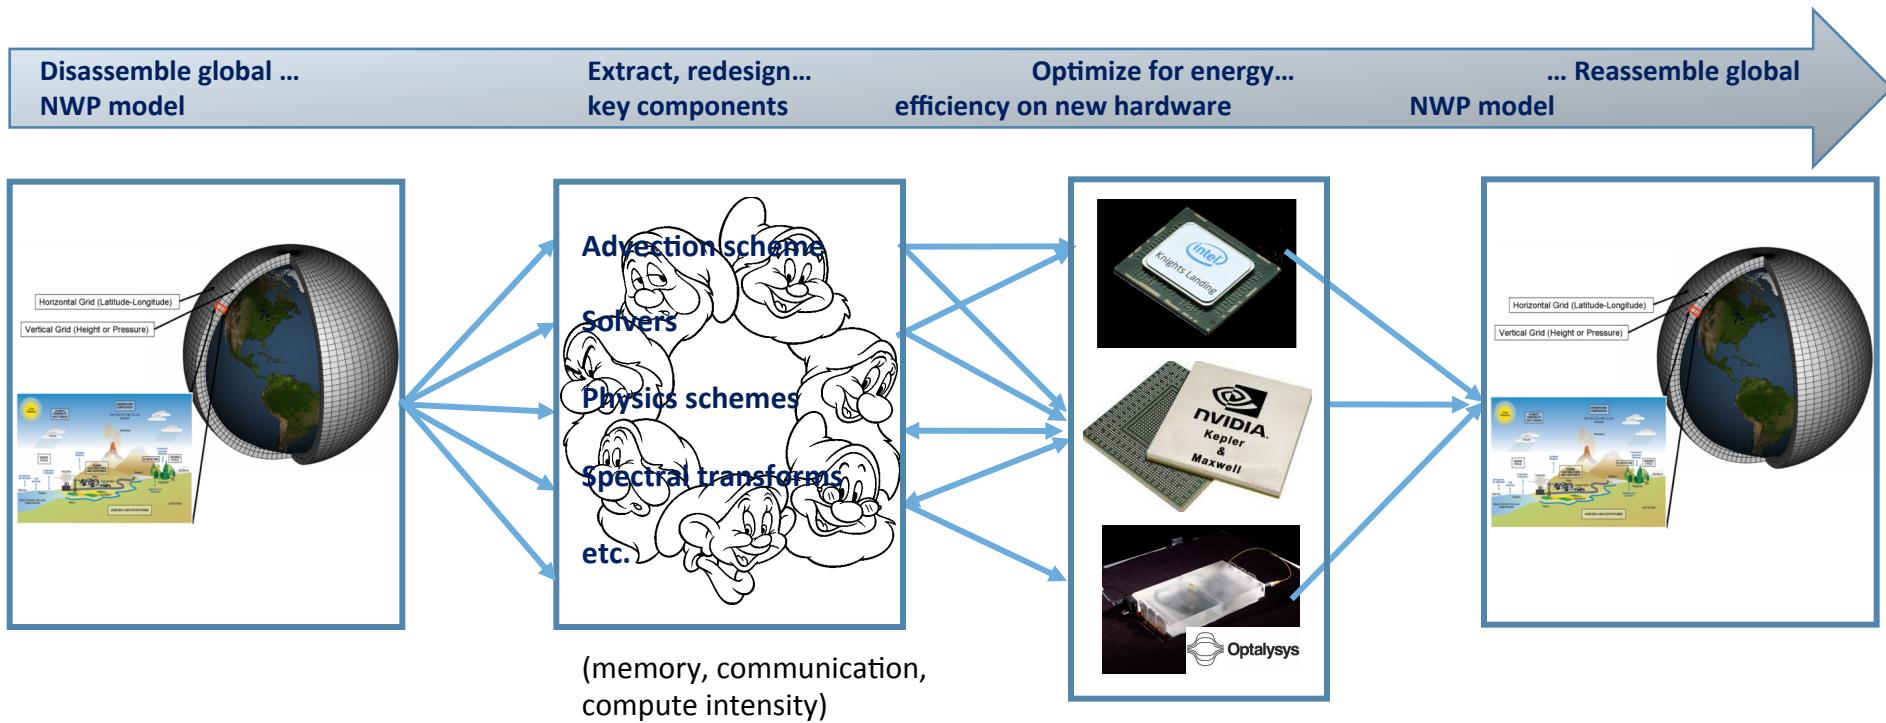
# ESCAPE

Disassemble global ...  
NWP model

Extract, redesign...  
key components

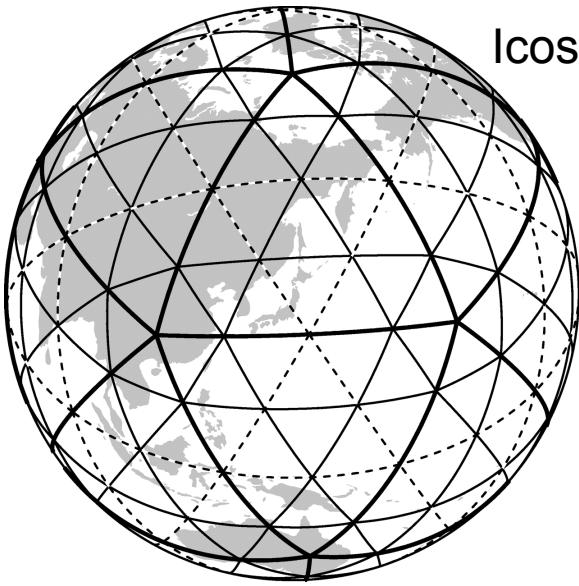
Optimize for energy...  
efficiency on new hardware

... Reassemble global  
NWP model

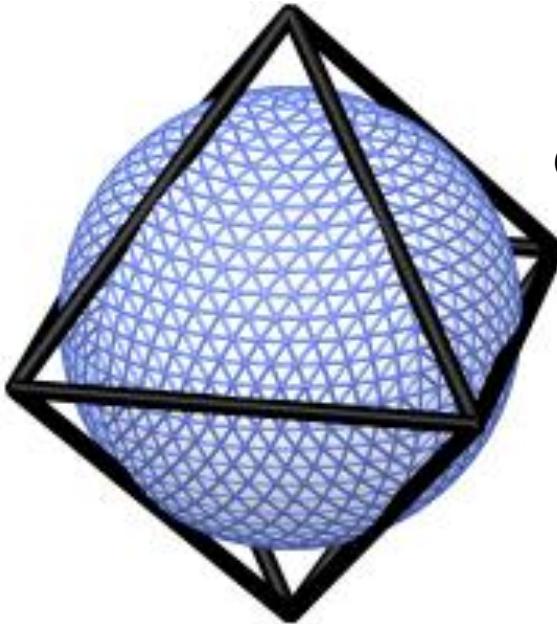




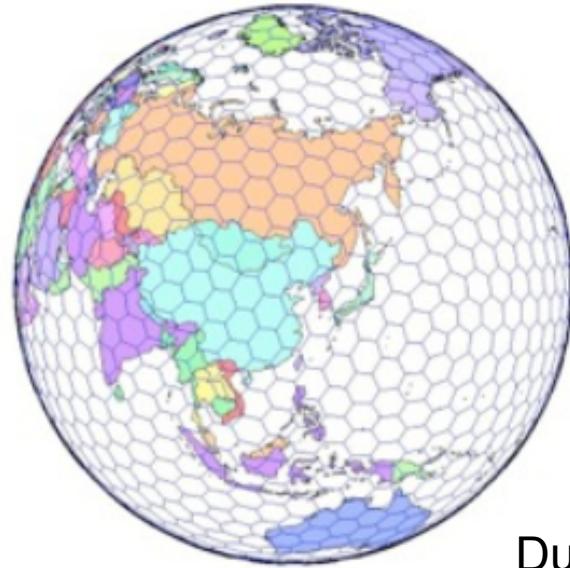
# GridTools On Irregular Grids



Icosahedral



Octahedral



Dual-grid

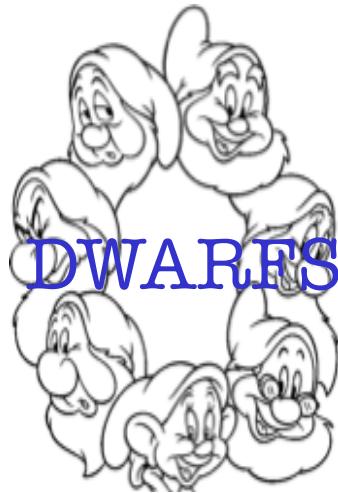


Cube sphere



# Dynamics

Advection (SL, MPDATA)  
 Elliptic solver (GCR)  
 Spectral transform (SH+)



**DWARFS**

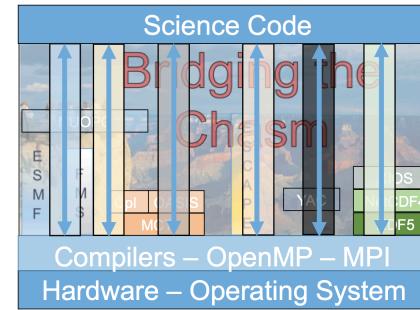
# Physical parameterisation

Cloud Microphysics  
 (CloudSC)  
 Radiation  
 (ACRANEB2)

Dwarf	Priority	Target Accelerators			Programming Models			
		GPU	Xeon Phi	Optalysis	MPI	OpenMP	OpenACC	DSL
D - spectral transform - SH	I	✓	✓	✗	✓	✓	-pencil	✗
D - spectral transform - biFFT	I	✓	✓	✓	✓	✓	-pencil	✗
D - advection - MPDATA	I	✓	✓	✗	✓	✓	-pencil	-pencil
I - LAITRI (3d interpol. algorithm)	I	✓	✓	✗	✗	✓	-pencil	-pencil
D - elliptic solver - GCR	II	✓	✓	✗	✓	✓	-pencil	sunburst
D - advection - semi-Lagrangian	II	✓	✓	✗	✓	✓	-pencil	-pencil
P - cloud microphysics - CloudSC	II	✓	✓	✗	✗	✓	sunburst	✗
P - radiation scheme - ACRANEB2	II	✓	✓	✗	✗	✓	-pencil	✗



From the bottom:

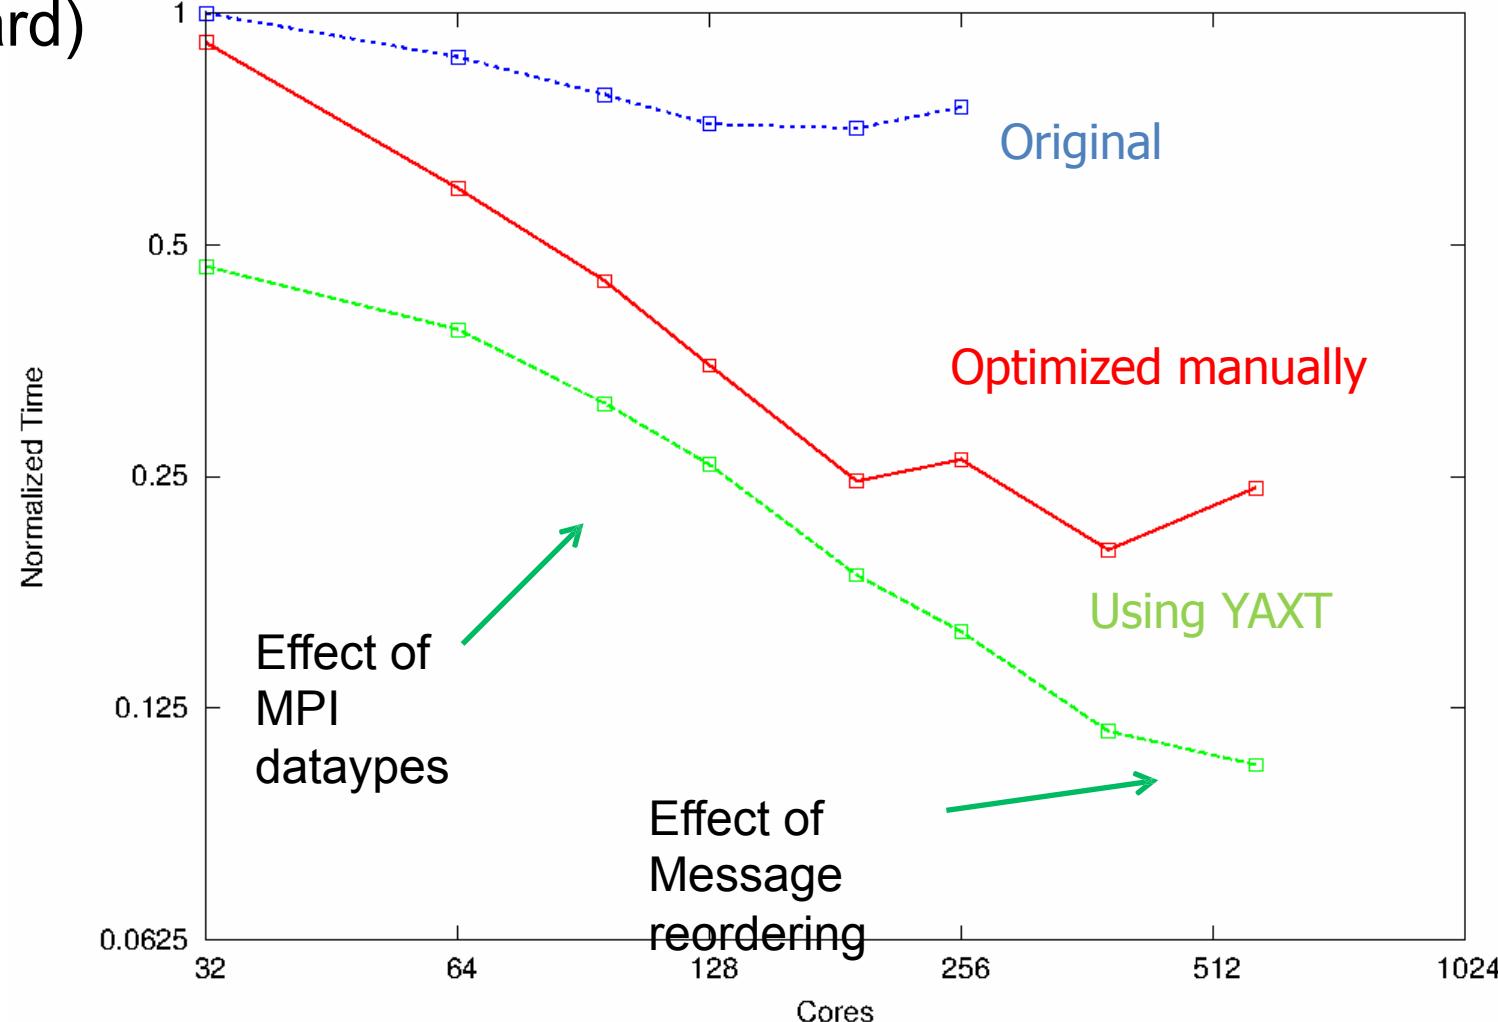


# WORK SMARTER WITH MPI

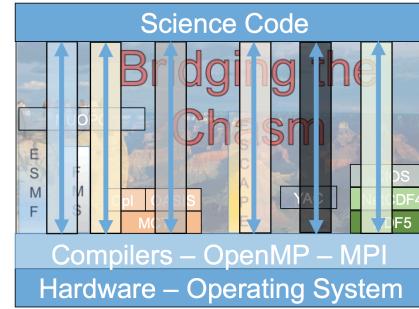
# YAXT: performance – example ECHAM

Global Transposition gp->ffsl

T63L47 (synchronized measurement on blizzard)



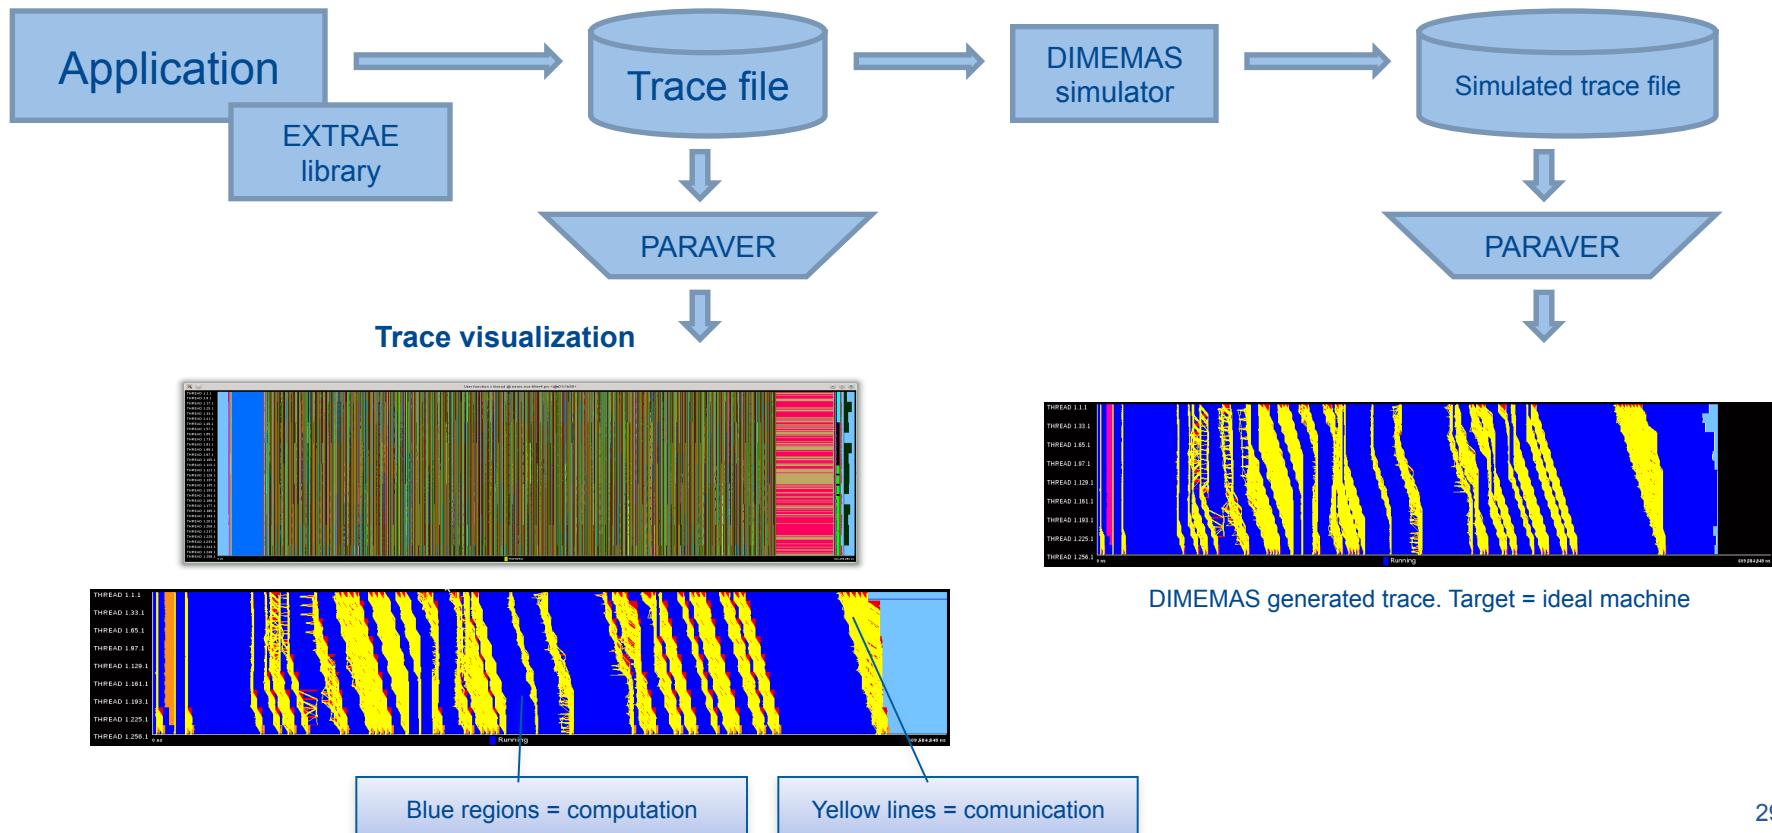
From the bottom:



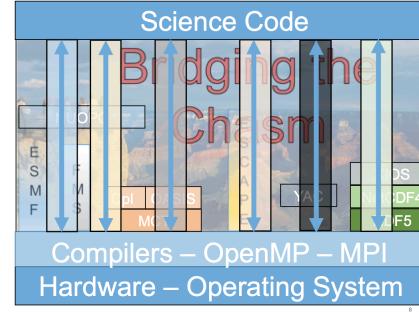
# UNDERSTAND WHAT IS GOING ON

## BSC-CNS performance tools:

- EXTRAE: Trace generation package.
- PARAVER: Trace visualization and analysis tool.
- DIMEMAS: Simulation tool for analysis on a configurable target platform.



IT'S A PEOPLE THING



# WHAT MAKES A SUCCESSFUL COMMUNITY TOOL?

Talking is not enough!

We need to work better together!

Can we take just one small step  
forward together?

What are the processes and  
procedures?

1. Nothing has to be used by everybody
2. What makes something a community code? OASIS Example
  - i. It helps to be the only game in town.
  - ii. Not very disruptive of existing solutions. Cost of transition/implementation.
  - iii. User support (tools for collaborative development, responsiveness a cost, who pays?) Training. Documentation. (We all support our own tools, but hide the cost?)
  - iv. Integrate developments back from the user community (another cost, what benefit?)
  - v. Momentum (in one direction, what about evolution?)
  - vi. People don't need to see development path if things are relatively satisfactory?
  - vii. Reward structures for new innovation in someone else's package?
  - viii. Courage to stop, and pick up and use/integrate others?
  - ix. Cf MPIM new activity: YAC (why not reintegrate?)

# Process?

1. Open Source
2. Open Development
3. Understanding of community/commitment/timescale?

Separation of concerns between objectives and implementations? What body exists or could be created that could talk about what to do, as opposed to how?

Ongoing  
role for  
ENES – the community,  
and  
ESIWACE – (a) project