

# Software stack deployment for Earth System Modelling

Sergey Kosukhin  
*Max Planck Institute for Meteorology*

IS-ENES2 Workshop on Workflow Solutions and Metadata Generation

29 September 2016  
Lisbon, Portugal



# Different users want different things

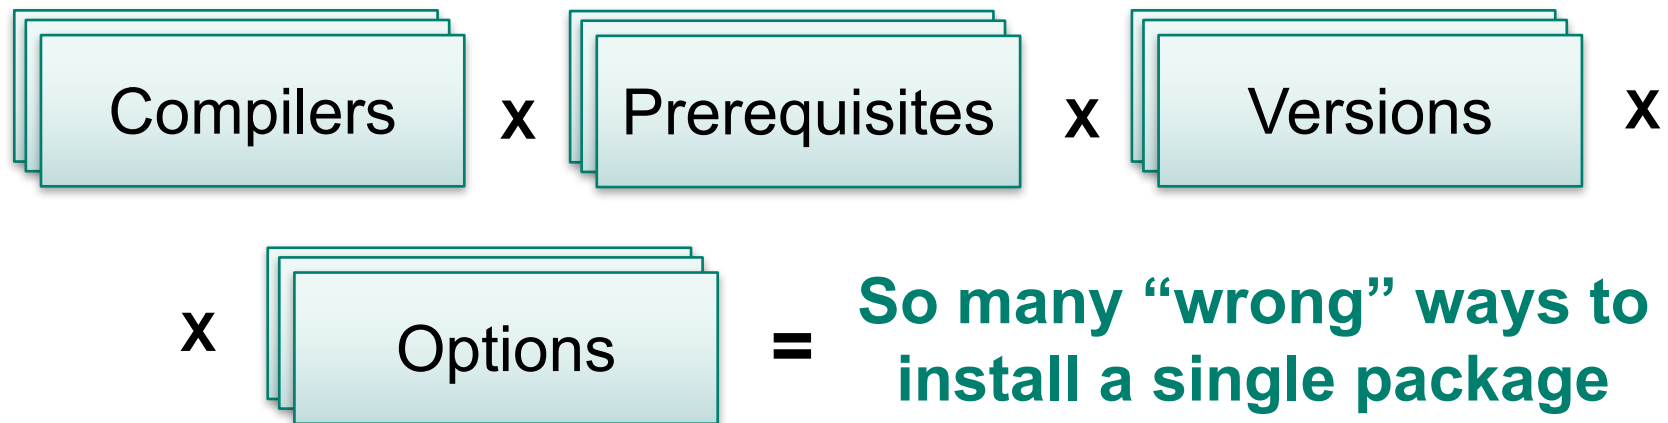
- Single installation of a particular model version (summer school).
- Several installations of the same model with different compilers, libraries, their versions and options (benchmarking/profiling).
- Continuous changes of source code with preferred toolchain (development).

*Keep installation process simple and flexible: the less users now about software, the less decisions they want to make and vice versa.*

# Environments are very different

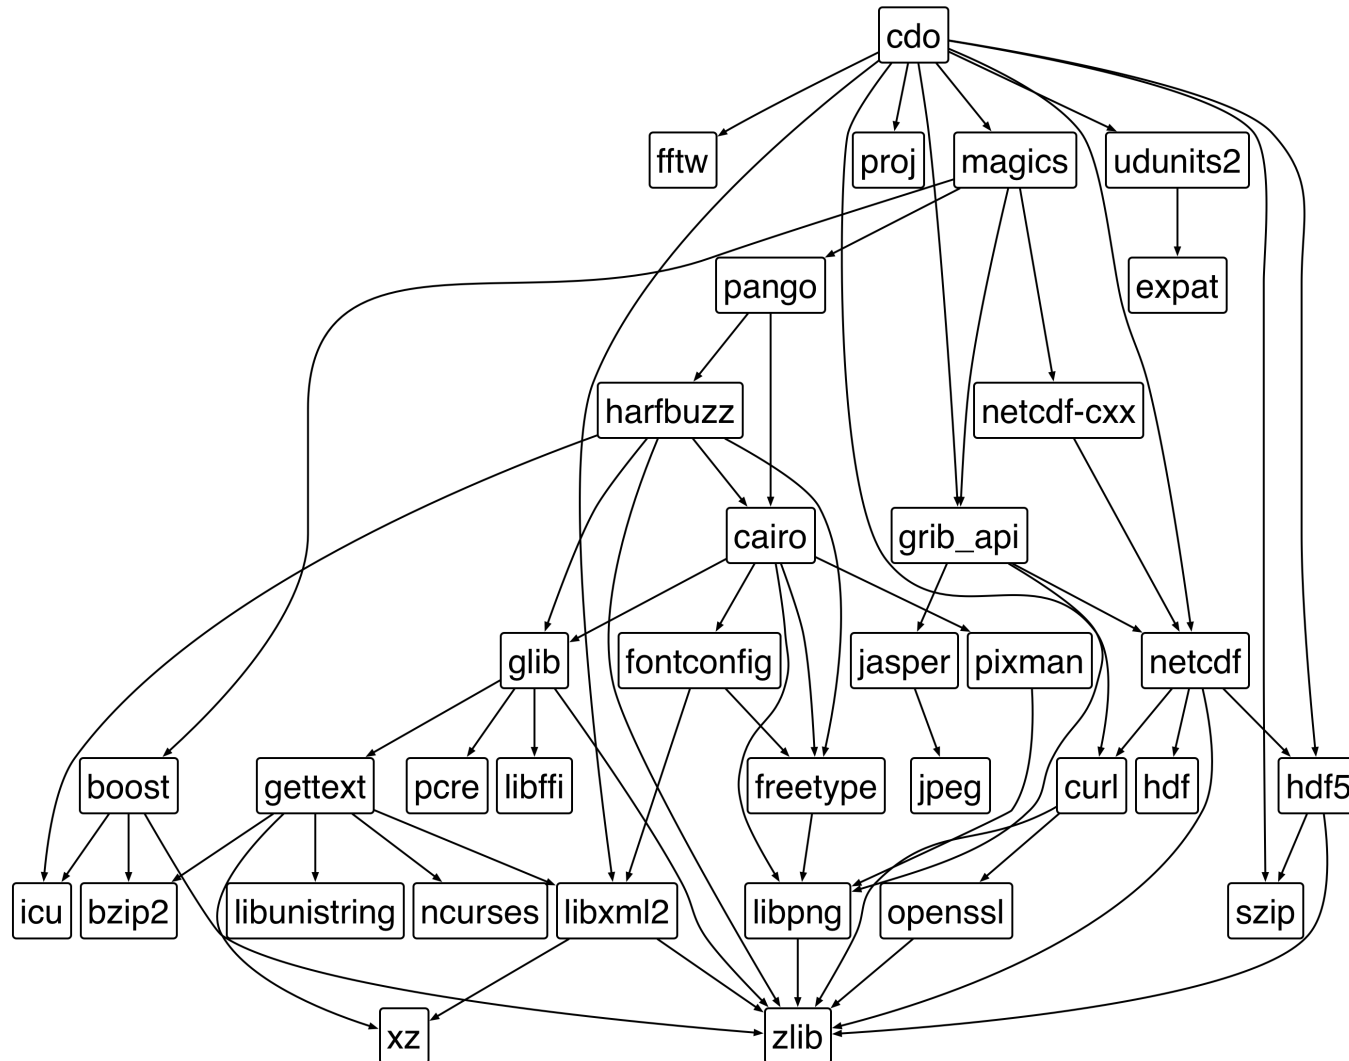
- Single machine or large-scale HPC site?
- Build everything from scratch or use provided system software?
- Which compiler? Which prerequisite packages and their versions?

*There isn't a standard way to build!*




*We don't always know in advance which one is right!*

# Dependency tree of the CDOs



# Existing tools

- Binary package managers
  - Designed to manage a single, stable and well tested stack.
  - Install one version of each package in a single prefix (/usr).
- Port systems
  - Macports, Homebrew, Gentoo, etc.
  - Minimal support for builds parameterized by compilers, dependency versions.
- Virtual Machines and Linux Containers (Docker)
  - Containers allow users to build environments for different applications.
  - Does not solve the build problem (someone has to build the image)
  - Performance, security, and upgrade issues prevent widespread HPC deployment.



# Spack is a flexible package manager



Lawrence Livermore National Laboratory

- How to install Spack:

**Get from git repository:**

```
$ git clone https://github.com/LLNL/spack.git
```

**Or download the archive and unzip it:**

```
$ wget https://github.com/LLNL/spack/archive/develop.zip
```

```
$ unzip develop.zip
```

**Setup environmental variables:**

```
$ . ./spack/share/spack/setup-env.sh
```


- How to install a package:

```
$ spack install hdf5
```

*Spack will detect available compilers and install HDF5 with all its dependencies.*

# Combinatorial complexity

Dependency DAG



Hash

Installation Layout

```
spack/opt/  
linux-debian7-x86_64/  
gcc-4.9.2/  
netcdf-4.4.1-dthu7lv5ml6g/  
netcdf-4.4.1-z33mrewv6iqd/  
hdf5-1.10.0-patch1-4zay7kzl2kcd/  
mpich-3.2-76exppmhc5z7/  
openmpi-2.0.0-pq5qkqmaeznk/  
...  
intel-16.0.2/  
netcdf-4.4.1-6wbb15ivdylx/  
hdf5-1.8.15-lkf14aq3nqiz/  
openmpi-2.0.0-3taka3nrbwwaq/  
...
```

- Each unique dependency graph is a unique configuration.
- Each configuration installed in a unique directory:
  - Configurations of the same package can coexist.
- **Hash** of entire directed acyclic graph (DAG) is appended to each prefix.
- Installed packages automatically find dependencies:
  - Spack embeds RPATHs in binaries;
  - No need to use modules or set LD\_LIBRARY\_PATH.


# Spack provides a spec syntax to describe customized DAG configurations

\$ spack install cdo	unconstrained
\$ spack install cdo@1.7.2	@ custom version
\$ spack install cdo@1.7.2 %gcc@4.9.2	% custom compiler
\$ spack install cdo@1.7.2 %gcc@4.9.2 +grib_api	+/~ build option
\$ spack install cdo@1.7.2 os=SuSE11	os=<frontend OS>
\$ spack install cdo@1.7.2 os=CNL10	os=<backend OS>
\$ spack install cdo@1.7.2 os=CNL10 target=haswell	target=<cpu target>

- Each expression is a **spec** for a particular configuration
  - Each clause adds a constraint to the spec
  - Constraints are optional – specify only what you need.
  - Customize install on the command line!
- Syntax abstracts details in the common case
  - Makes parameterization by version, compiler, and options easy when necessary




# Spack specs can constrain versions of dependencies



```
$ spack install netcdf %intel@16.0.2 ^zlib@1.2.8
```

- Spack ensures one configuration of each library per DAG
- Spack can ensure that builds use the same compiler

# Spack handles ABI-incompatible, versioned interfaces like MPI



- *mpi is a virtual dependency*
- Install the same package built with two different MPI implementations:

```
$ spack install netcdf ^mvapich@1.9
```

```
$ spack install netcdf ^openmpi@1.4:
```

- Let Spack choose MPI version, as long as it provides MPI 2 interface:

```
$ spack install netcdf ^mpi@2
```

# Spacks allows optional dependencies

- The user can define named variants (flags):

```
variant("python", default=False, "Build with python support")  
depends_on("python", when="+python")
```

- And use them to install:

```
$ spack install vim +python  
$ spack install vim -python
```

- Dependencies may be optional according to other conditions:

e.g., gcc dependency on mpc from 4.5 on:

```
depends_on("mpc", when="@4.5:")
```

# Spack packages are Python scripts

```
from spack import *

class Magics(Package):
    homepage = "https://software.ecmwf.int/wiki/display/MAGP/Magics"
    url = "https://software.ecmwf.int/wiki/download/attachments/3473464/Magics-2.29.0-Source.tar.gz"

    version('2.29.4', '91c561f413316fb665b3bb563f3878d1')
    version('2.29.0', 'db20a4d3c51a2da5657c31ae3de59709', preferred=True)

    patch('no_hardcoded_python.patch')
    patch('resolve_isnan_ambiguity.patch', when='@2.29.0')

    variant('bufr', default=False, description='Enable BUFR support')
    # More variants here...

    depends_on('grib_api')
    # More dependencies here...
    depends_on('libemos', when='+bufr')

    def install(self, spec, prefix):
        options = []
        options.extend(std_cmake_args)
        options.append('-DGRIB_API_PATH=%s' % spec['grib_api'].prefix)

        if '+bufr' in spec:
            options.append('-DENABLE_BUFR=ON')
            options.append('-DLIBEMOS_PATH=%s' % spec['libemos'].prefix)
        else:
            options.append('-DENABLE_BUFR=OFF')

        with working_dir('spack-build', create=True):
            cmake('..', *options)
            make()
            make('install')
```

Metadata

Versions

Patches

Variants

Dependencies

Commands for installation


# Concretization fills in missing configuration details when the user is not explicit.

mpileaks ^callpath@1.0+debug ^libelf@0.8.11

User input: *abstract* spec with some constraints

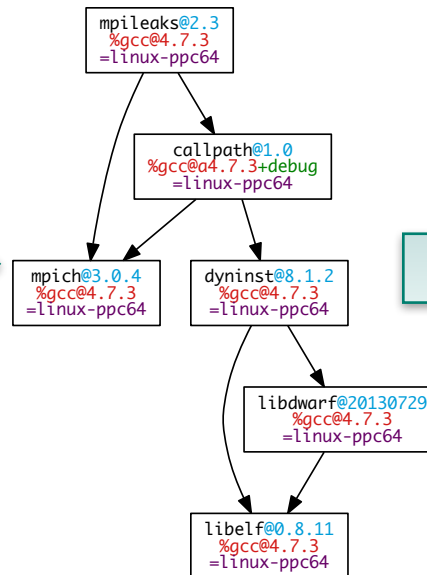
spec.yaml

Normalize



*Abstract*, normalized spec with some dependencies.

Concretize



*Concrete* spec is fully constrained and can be passed to install.

Store

```
spec:
- mpileaks:
  arch: linux-x86_64
  compiler:
    name: gcc
    version: 4.9.2
  dependencies:
    adept-utils: kszrtkpbzac3ss2ixcjkcorlaybnpt4
    callpath: bah5f4h4d2n47mgycej2mtrnrivvxy77
    mpich: aa4ar6ifj23yijqmdabeakpejcli72t3
    hash: 33hjhxii7p6gyzn5ptgyes7sghyprujh
    variants: {}
    version: '1.0'
- adept-utils:
  arch: linux-x86_64
  compiler:
    name: gcc
    version: 4.9.2
  dependencies:
    boost: teesjv7ehpe5ksspjim5dk43a7qnowlq
    mpich: aa4ar6ifj23yijqmdabeakpejcli72t3
    hash: kszrtkpbzac3ss2ixcjkcorlaybnpt4
    variants: {}
    version: 1.0.1
- boost:
  arch: linux-x86_64
  compiler:
    name: gcc
    version: 4.9.2
  dependencies: {}
  hash: teesjv7ehpe5ksspjim5dk43a7qnowlq
  variants: {}
  version: 1.59.0
...
```

Detailed provenance is stored with the installed package

# `spack find` shows what is installed

```
$ spack find
==> 54 installed packages.
-- linux-x86_64 / gcc@4.4.7 -----
ImageMagick@6.8.9-10  glib@2.42.1      libtiff@4.0.3
SAMRAI@3.9.1         graphlib@2.0.0    libtool@2.4.2
adept-utils@1.0      gtkplus@2.24.25   libxcb@1.11
atk@2.14.0           harfbuzz@0.9.37   libxml2@2.9.2
boost@1.55.0         hdf5@1.8.13       llvm@3.0
cairo@1.14.0         icu@54.1          metis@5.1.0
callpath@1.0.2       jpeg@9a           mpich@3.0.4
dyninst@8.1.2        libdwarf@20130729 ncurses@5.9
dyninst@8.1.2        libelf@0.8.13     ocr@2015-02-16
fontconfig@2.11.1    libffi@3.1        openssl@1.0.1h
freetype@2.5.3       libpng@2.0.2      otf@1.12.5salmon
gdk-pixbuf@2.31.2    libpng@1.6.16     otf2@1.4

-- linux-x86_64 / gcc@4.8.2 -----
adept-utils@1.0.1    boost@1.55.0      cmake@5.6-special
adept-utils@1.0.1    cmake@5.6         dyninst@8.1.2

-- linux-x86_64 / intel@14.0.2 -----
hwloc@1.9  mpich@3.0.4  starpu@1.1.4

-- linux-x86_64 / intel@15.0.0 -----
adept-utils@1.0.1  boost@1.55.0  libdwarf@20130729

-- linux-x86_64 / intel@15.0.1 -----
adept-utils@1.0.1  callpath@1.0.2  libdwarf@20130729
boost@1.55.0       hwloc@1.9       libelf@0.8.13
```

- All the versions coexist!  
Multiple versions of same package are ok.
- Packages are installed to automatically find correct dependencies.
- Binaries work regardless of user's environment.
- Spack also generates module files.  
Don't have to use them.

# Adding compiler flags (for compiler parameter studies)

```
$ spack install ncl cflags=\'-03 -g -fast -fpack-struct\'
```

- This would install NCL with the specified flags
  - Flags are injected via Spack's compiler wrappers.
- Flags are propagated to dependencies automatically
  - Flags are included in the **DAG hash**
  - Each build is considered a **different version**
- This provides an easy harness for doing parameter studies for tuning codes

# Status of the integration of the ESM applications into Spack

- CDO (<https://code.zmaw.de/projects/cdo>): fully featured, main branch
- Magics (<https://software.ecmwf.int/wiki/display/MAGP/Magics>): almost fully featured (no python support yet), main branch
- Emoslib (<https://software.ecmwf.int/wiki/display/EMOS/Emoslib>): minimum implementation, main branch
- NCAR NCL (<https://www.ncl.ucar.edu>): WIP

## Plans on ES models integration:

- MPI-ESM1   • ICON
- EC-EARTH   • FAMOUS

This work is supported by the ESiWACE project funded by the European Union's Horizon 2020 Research and Innovation Programme



# Known issues

- Sometimes spack fails to prevent packages from linking against external libraries.
- Not all the packages are at the production level.
- Spack is under the continuous development without explicit stable version.
- URL lists of packages require maintenance.

# Conclusion

- There is no magic: the more standard the original building system of a package is the easier it is to integrate it into spack and vice versa.
- Among other things it is a good way to document the installation process.

# References

- T. Gamblin, M. LeGendre, M. R. Collette, G. L. Lee, A. Moody, B. R. de Supinski, and S. Futral. The Spack package manager: Bringing order to HPC software chaos. In *International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 40:1–40:12, 2015.
- T. Gamblin et al. The Spack package manager: Bringing order to HPC software chaos [PDF slides]. Retrieved from <https://tgamblin.github.io/files/Gamblin-Spack-SC15-Talk.pdf>
- T. Gamblin. Spack: a tool for managing HPC software [PowerPoint slides]. Personal correspondence.
- Spack documentation. <http://spack.readthedocs.io/en/latest>

# Thank you!

Sergey Kosukhin

*[sergey.kosukhin@mpimet.mpg.de](mailto:sergey.kosukhin@mpimet.mpg.de)*