



GFDL FMS

Bronx to Chaco Rewrite

IS-ENES WWMG 2016

Presented by Chan Wilson

29 September, 2016

Chan Wilson, Engility

Erik Mason, Engility

Chris Blanton, Engility

Karen Paffendorf, Princeton

Seth Underwood, US Federal

Jeff Durachta, Engility

V. Balaji, Princeton





In memoriam



Amy Langenhorst 1977-2016





Overview

- Intro to GFDL, FMS, FRE
- Climate modeling workflow achievements
- How we're approaching a workflow rewrite
- Challenges, lessons learned



GFDL

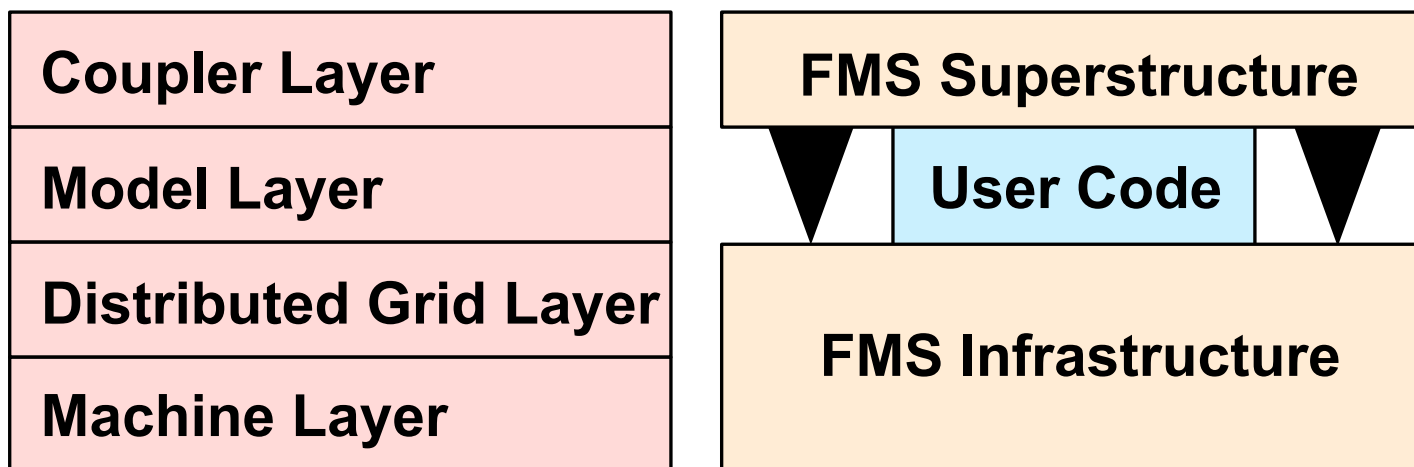
Geophysical Fluid Dynamics Laboratory

- Mission: *To advance scientific understanding of climate and its natural and anthropogenic variations and impacts, and improve NOAA's predictive capabilities*
- Joint Institute with Princeton University
- About 300 people organized in groups and teams
 - 8 scientific groups
 - Technical Services
 - Modeling Services
 - Framework for coupled climate modeling (FMS)
 - Workflow software (FRE), Curator Database
 - Liaisons to scientific modeling efforts



Flexible Modeling System

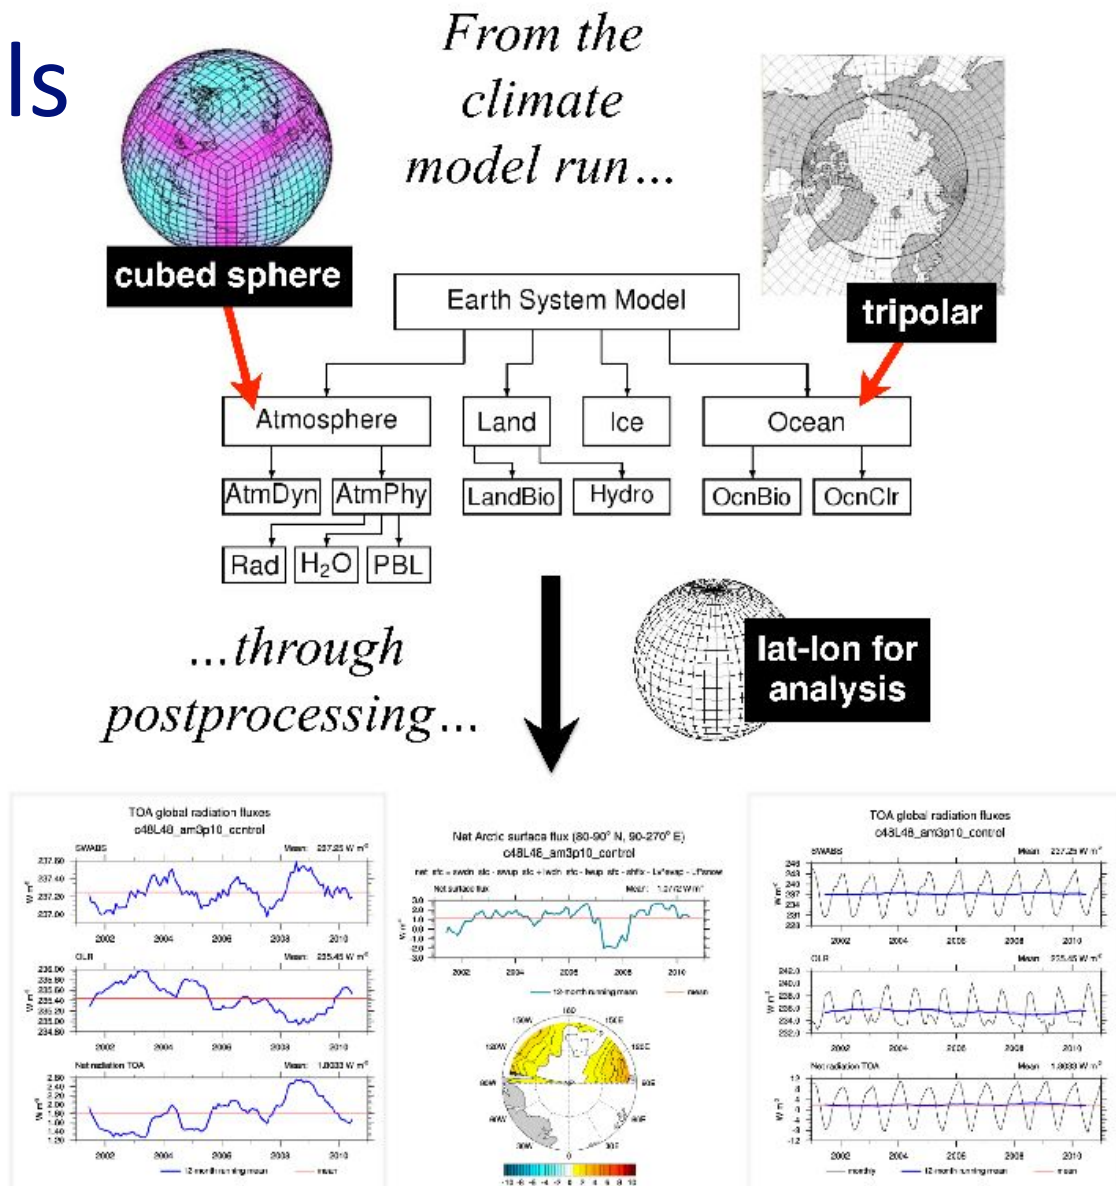
- FMS is a software framework which provides infrastructure and interfaces to component models and multi-component models.



- Started ~1998, development timelines:
 - Component models by scientific group, continuous
 - Annual FMS city releases, 200+ configurations tested

Workflow Goals

- Reproducibility
- Perturbations
- Differencing
- Robustness
- Efficiency
- Error handling:
user level &
system level

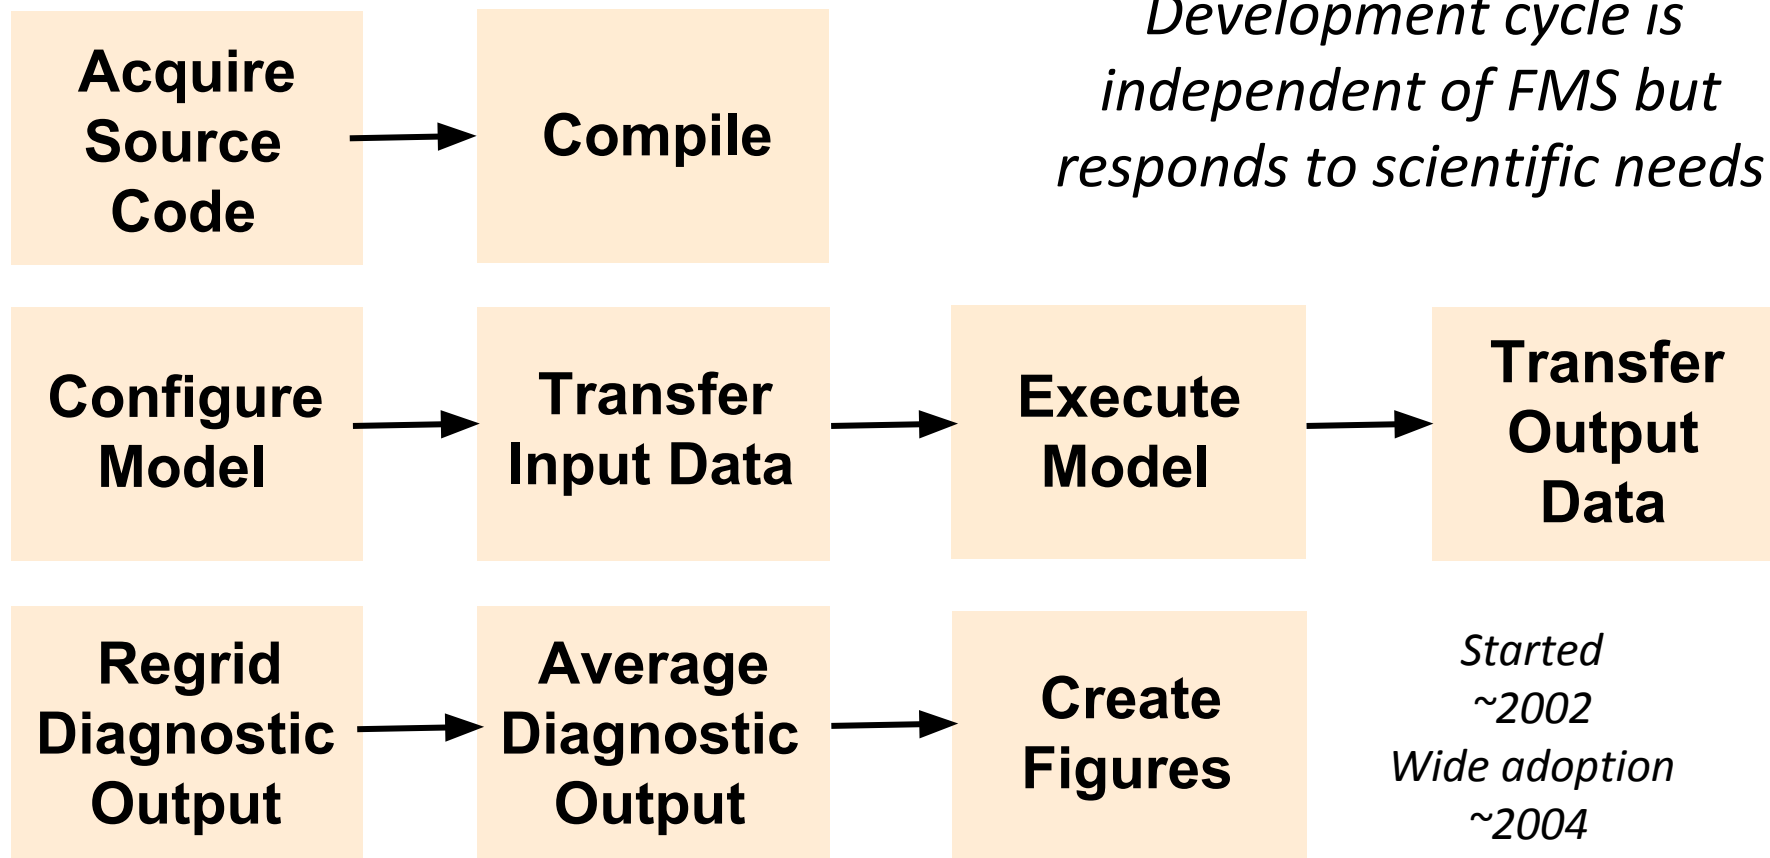


...to automatically generated figures.



FMS Runtime Environment

FRE is a set of workflow management tools that read a model configuration file and take various actions.



FRE Technologies

- XML model configuration files, and schema
 - User and canonical XML
- Perl command line utilities read XML file
- Site configuration files
 - Script templates
 - Conventions for file names and locations
- Environment modules to manage FRE versions
- Moab / Torque to schedule jobs across sites
- Transfer tools: globus + local gcp, HSM tools
- NCO Operators + local fregrid, plevel, timavg
- Jenkins for automated testing
- *Coming soon: Cylc*

FRE Features

- Encapsulated support for multiple sites
 - Compiler, scheduler, paths, user defaults
- Model component-based organization of model configuration
- Integrated bitwise model testing
 - Restarts, scaling (vs. reference and self)
- Experiment inheritance for perturbations
- User code plug-ins
- Ensemble support
- Postprocessing and publishing, browsable curator database




FRE Job Stream Overview

- Experiments run in segments of model time
 - Segment length is user-defined
 - More than one segment per main compute job is possible
- After each segment:
 - State, restart, diagnostic data is saved
 - Model can be restarted (checkpointing)
 - Data is transferred to long term storage
 - A number of post-processing jobs may be launched, highly task parallel





Flow Through The Hardware



1. Remote transfer of input data
- 2a. Transfer/preprocess input data
- 2b. Model execution
- 2c. Transfer/postprocess output data
3. Remote transfer of output data
4. Post-processing

GFDL

Image:
Tara McQueen

11 / 26





GFDL Data Stats

- Networking capacity ORNL to GFDL: two 10gig pipes, 120 TB/day theoretical for each pipe
- Analysis cluster:
 - ~100 hosts with 8 to 16 cores, 48GB to 512GB memory, local fast scratch disk of 9TB to 44TB
 - 4 PB/week throughput
- Tape-based data archive: 60PB
 - ~2PB disk attached
 - Another ~2PB filesystem shared among hosts for caching intermediate results
- 1300 auto-generated figures from atmos model
- An early configuration of CM2.6 (0.1 degree ocean, 0.5 degree atmos) runs on 19,000 cores
 - a year of simulation takes 14 hours and generates 2TB of data per simulation year, ran 300 simulation years

Post-processing Defined

- Preparing diagnostic output for analysis
 - Time series: Hourly, daily, monthly, seasonal, annual
 - Annual and Seasonal climatological averages
 - Horizontal interpolation: data on various grids can be regridded to lat-lon with “fregid”
 - Vertical interpolation: to standard pressure levels
 - Hooks to call user scripts to create plots or perform further data manipulation
 - Enter the model into the curator database
- Requirement: must run faster than the model
- Self healing workflow: state is stored; tasks know their dependencies and resubmit as needed



10 Years Later...

Oh, the clusters you'll run on and the lessons you'll learn...



FMS FRE Chaco

- Rewrite of FRE beginning with post processing
- Maintain compatibility and historical behavior, yet standardize tool behavior
 - Old (user) interfaces available, e.g. ‘drop in’ replacement where possible
- Improve visibility and control of experiments

Chaco Major Goals

- Robustness and reliability
- Support for high resolution resource requirements
 - CM2.6, a higher resolution model, generates 2TB per simulation year, completes 2 sim years/day
- Support for discrete toolset that can be used without running end to end “production frepp”
- Monitoring
- Increased task parallelism
- Maintain existing functionality
- Keep pace with data flow rate from remote computing sites (gaea/theia/titan...)

High Resolution Strategy

- Reduce memory and disk space requirements
- Initially break all diagnostic data up into “shards”
 - “Shard” files contain one month’s worth of data for one variable
 - Shards can be on model levels or regridded
 - Perform data manipulations on one variable at a time, then combine data later if necessary
 - Operations on shards can be highly parallelized
- Make intelligent use of disk cache with intermediate data
- Reducing data movement is key
 - Overwhelming majority of the time spent in postprocessing is simply moving data



Black box to discrete toolset

Postprocessing

- Outside looks the same

```
frepp -v -c split -s -D -x FILE.XML -P  
SITE.PLATFORM-COMPILER -T TARGET-STYLE -d  
MODEL_DATA_DIRECTORY -t START_TIME EXPERIMENT_NAME
```

- Inside nothing is the same





FRE Utilities

Bronx

Chaco

fremake obtain code, create/submit compile scripts

frerun create and submit run scripts

frepp create and submit post-processing scripts

frelist list info about experiments in an xml file

frestatus show status of batch compiles and runs

frecheck compare regression test runs

frepriority change batch queue information

freppcheck report missing post-processing files

frescrub delete redundant post-processing files

fredb interact with Curator Database

fre cylc run

fre cylc prepare

fre list

fre monitor



Bronx Design

- Bronx implementation: monolithic, linear perl script re-invoking itself over subsequent segments of model data. Steps to produce desired products hardcoded in blocks of shell which are assembled and submitted to batch scheduler.



Chaco Design, 1

- Discrete tools in a unified code base: data movement, refinement, interpolation, timeseries and timeaverage
`fre data {get, split, zinterp, xyinterp, refineDiag, analysis, timeaverage, timeseries}`
- Cylc-based cycle for experiment duration and model segments
- Per-cycle tool runs

Chaco Cylc

- Cylc for the dependency and scheduling engine
- Cylc suites generated by FRE Chaco
 - Allows for different grouping of tasks based on arguments, XML, or other factors.
 - User accessible / modifiable, but 'hidden'
- Leverage Cylc task management and job submission
- Cylc GUI gives visibility and control



Chaco Design, 2

- Segments run in parallel:
 - Multiple MOAB jobs on separate nodes
- Tool commands
 - Utilize FRED to store parsed XML, diagfield info
 - Run concurrently over model components and model variables (shards)
 - Log all stdout, stderr, cpu, memory utilization, execution time

Chaco Tool Outline

- Modularize the monolithic workflow
- Standardize interface
 - ‘fre CATEGORY ACTION [OPTIONS]’
- Can operate independent of workflow
- On-disk data structure (FRED) stashes parsed experiment details
- File inventory and tool run databases
- Tools know input and output files



Chaco Tech



ENGILITY
Engineered to Make a Difference

- Perl OO via Moose and friends
 - Path to Perl 6
- CPAN all things
- Test driven development with continuous integration tactics.
- Strong source code practices, code documentation and project management tools to enable the team.

CPAN





FRE PostProcessing, 1

- **Get, Split, and Stage**
Retrieve data from storage, separate into variable shards, copy to shared temp filesystem.
- **Interpolate in Z, and XY**
Stage shards in, interpolate, stage out
- **Generate Timeseries**
Stage shards in, timeseries, store product
- **Generate Timeaverage**
Stage shards in, timeseries, store product



FRE PostProcessing, 2

Reality check!

- **Get, “RefineDiag”, Split, and Stage**
Retrieve data from storage, run user-specified scripts to “refine data”, separate into variable shards, copy to shared temp filesystem.
- **Interpolate in Z, and XY**
Stage shards in, interpolate, stage out
- **Generate Timeseries, timeaverage**
Stage shards in, timeseries / timeavg, store product
- **“Analysis”**
Run user-specified scripts to generate figures, etc



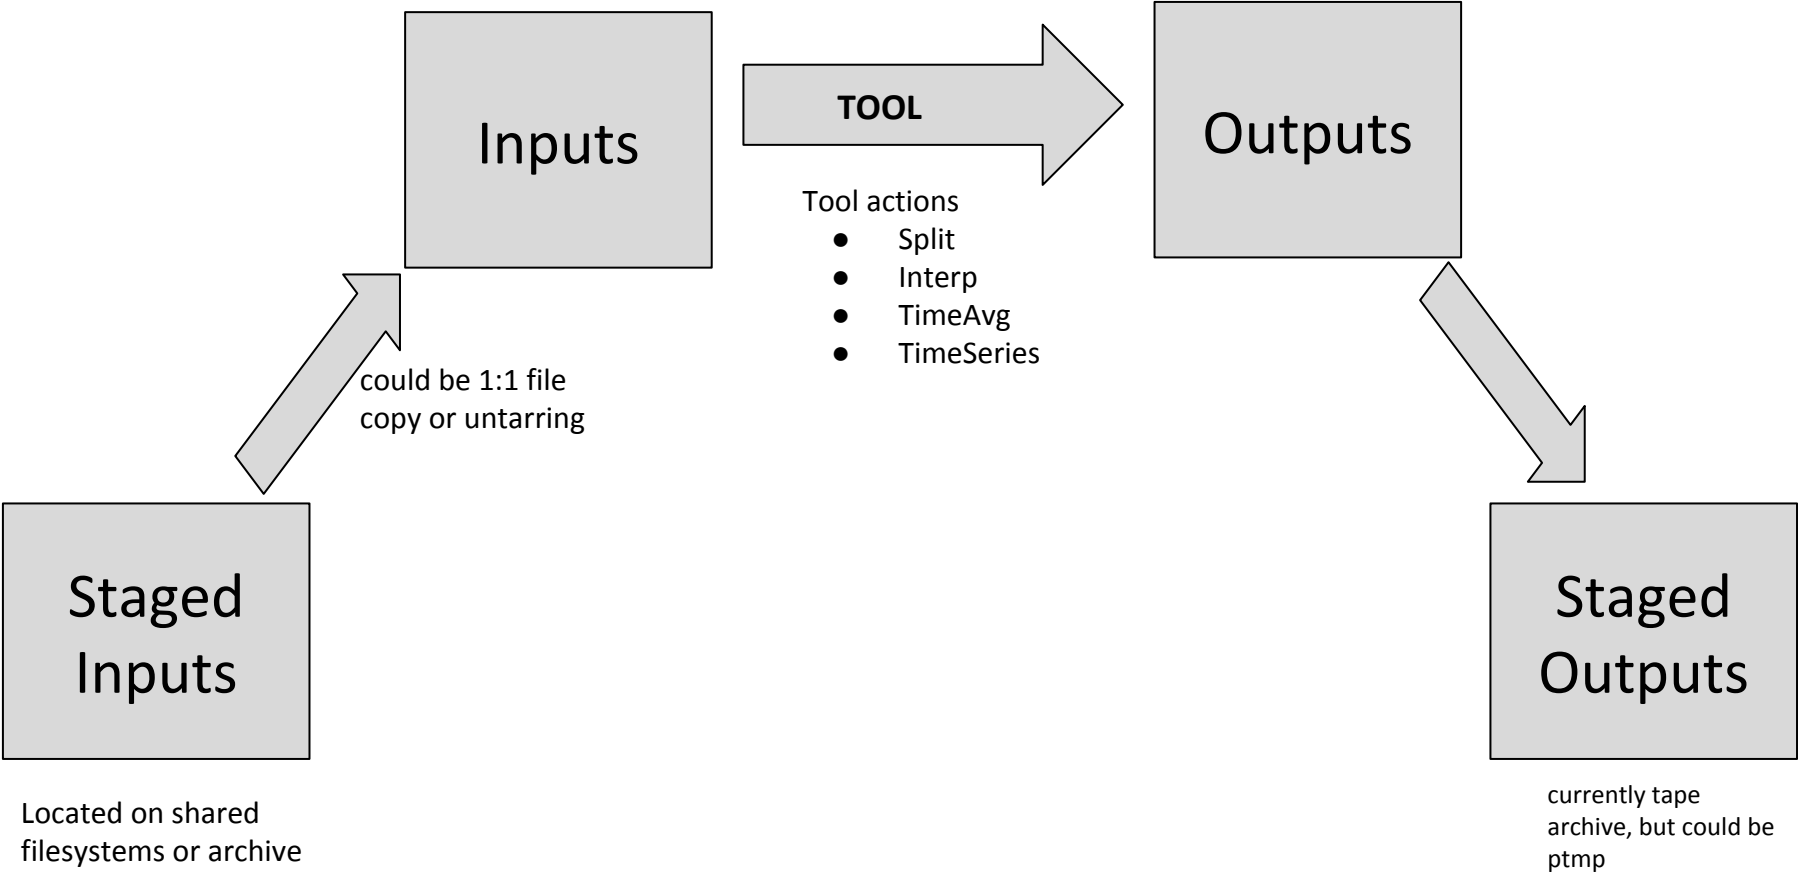


FRE PostProcessing, 3

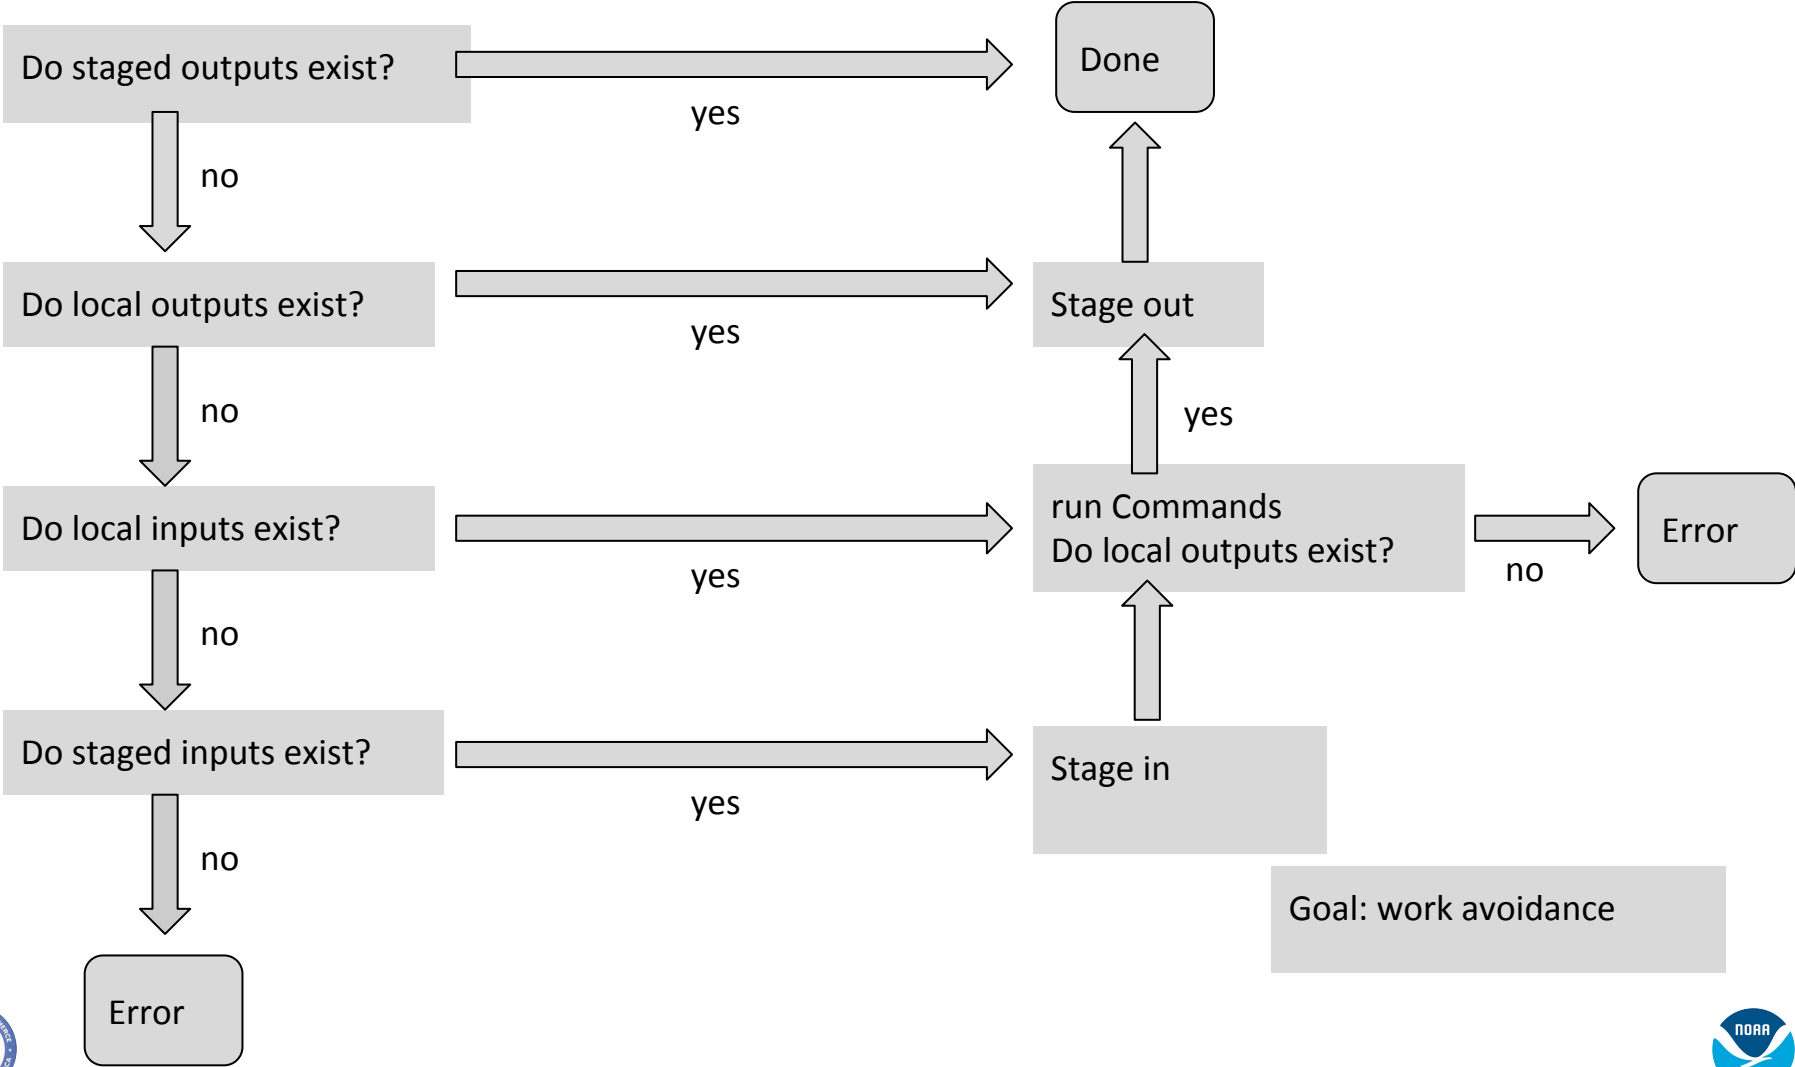
Dependency and Recovery check!

- Add knowledge of inputs and outputs to tools
- A Tool Operation framework for recovery and dependency-checking
 - Check for output (in final and local locations)
 - Check for inputs (in local and staged locations)
 - Run tool

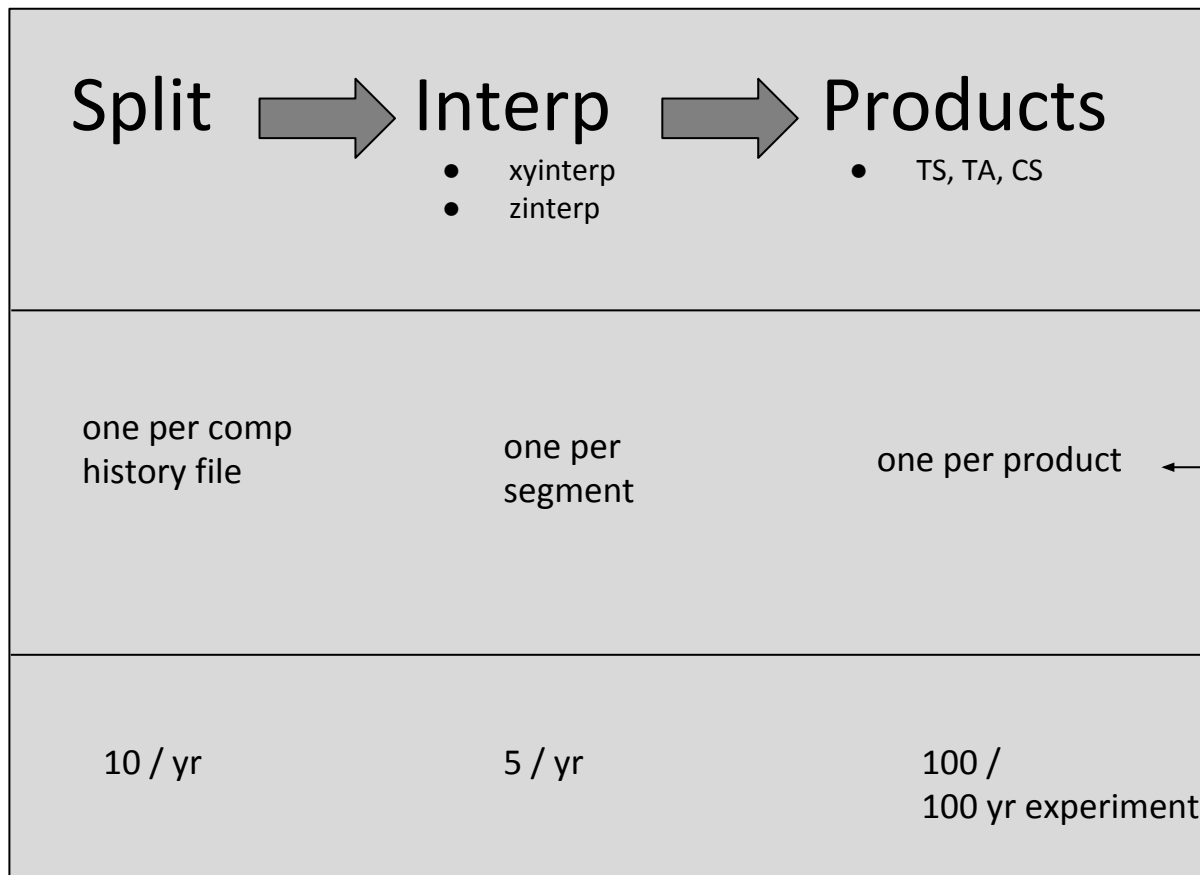
FRE Operation



Running an Operation



Operation Granularity



requests in XML (examples)

- 5 yr monthly TS ocean regridded to 1 deg
- All statics from land
- 10 yr seasonal TA atmos tracer native grid

Goals

- maximize outer-parallel
- make each be able to be a Cylc task without performance cost
- work avoidance

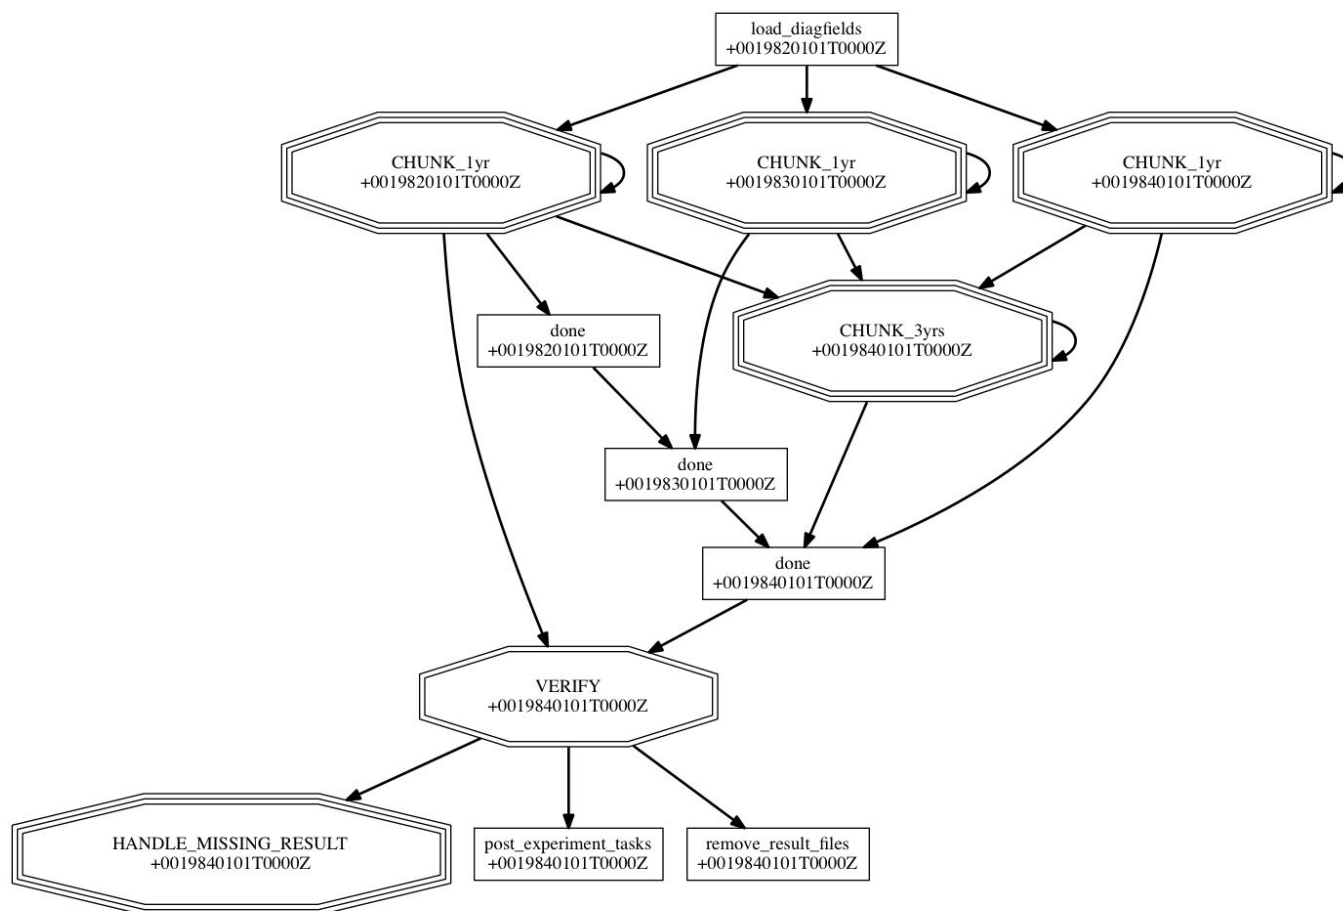
Tiers of concurrency

Goal: maximize inner-parallel by using tiers to isolate dependent commands

Example: One operation to create one annual timeaverage file containing all variables

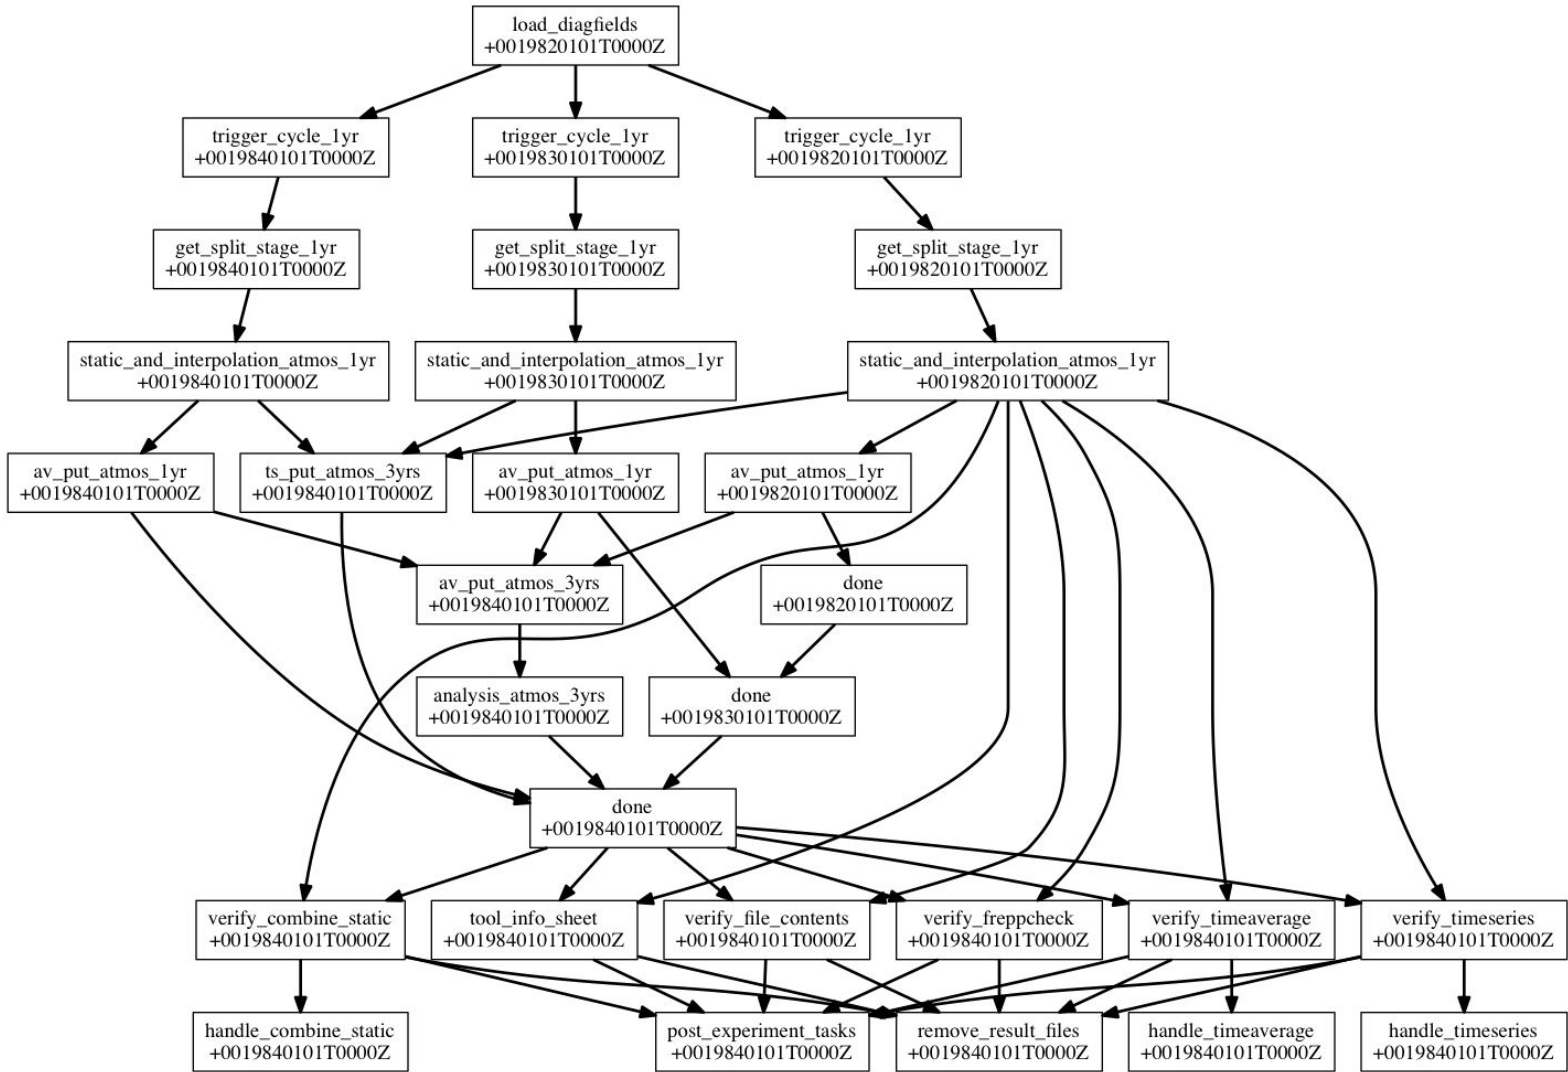
	Action	Tool	Parallelization
Tier 1	Concatenate shards to create TS	ncrcat	Number of variables
Tier 2	Average over TS timesteps	timavg (FRE tool)	Number of variables
Tier 3	Combine per-variable TA files into single file	ncks	1

3 year test and verify experiment





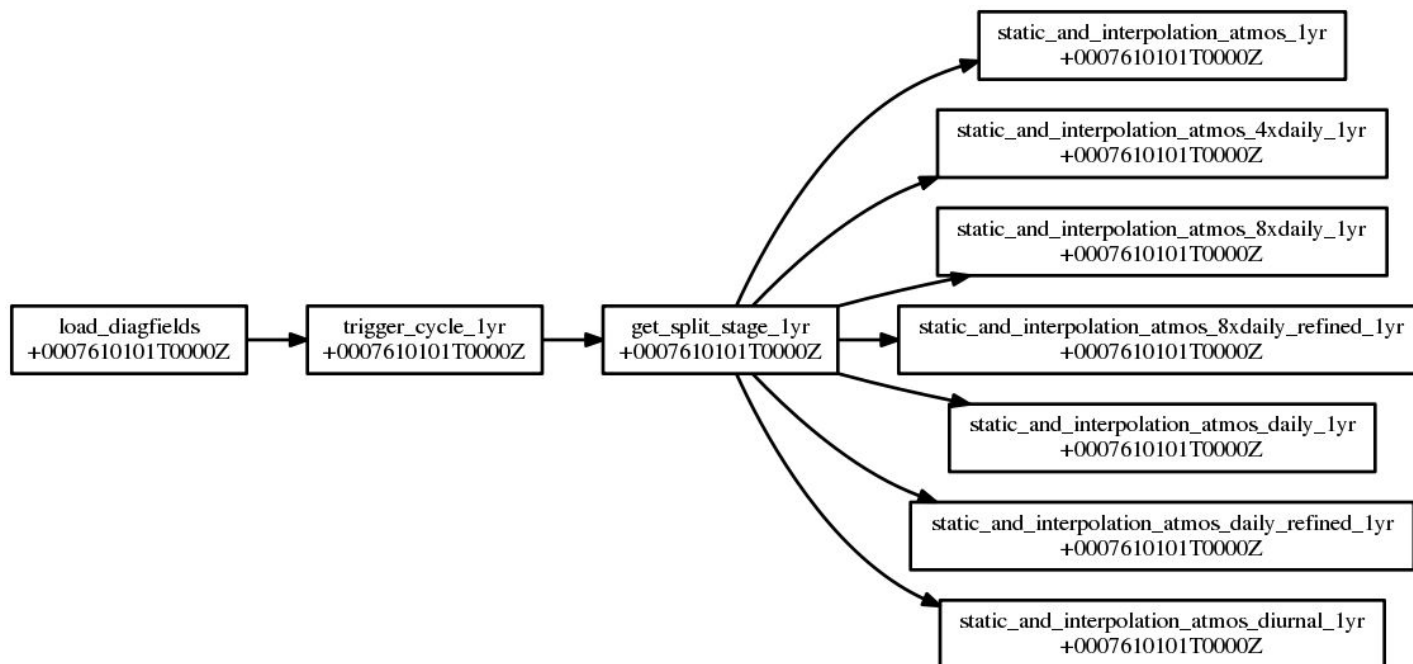
FRE Workflow, 3yr test



fre.chaco_testing_20160922T195204_98-quickstart.m45_am2p14_8thread.atmos.19820101-19840101



Simplified 1yr cycle of 100yr experiment



fre.verona.SM2_Control-1990.multiple.07610101-08600101

Current Status

- Alpha release available: *module load fre/chaco-alpha*
 - Have achieved full functionality for “quickstart” AM2, simplified c48 AM3, ESM2G, OM4, AWG
 - Differences in netcdf output vs. bronx exist but are understood
 - Workflow recovery dependency analysis
 - refineDiag: Create custom variables between run and postprocessing
 - Performance test performed; results currently under review
- Current work includes:
 - Adding ability to run in production from Bronx runscript
 - Review of performance data
 - Evaluate concurrency and parallelism
 - Adding experiment configurations to our test suite

Cool! So all the workflow issues
are solved! Right??



When is the workflow broken?

- As complexity increases, a smaller and smaller percentage of error is tolerable
 - If 1% of jobs fail and your workflow launches thousands of jobs, you still have lots of work
- Our strategy: standard logging, retries, recovery, error handling, Cylc, automated test infrastructure in Jenkins...



Analyzing Chaco 1

- Analyze timing and resource metrics

Performance Analysis of Chaco

- Tracking details at many levels
 - Individual commands (ncks, gcp)
 - Tiers of commands (run concurrently)
 - Command groupings
 - Segment runs
 - Tool runs



Analyzing Chaco 2

- What can we determine?
 - Variability
 - Data movement costs
 - Resource usage
- Areas under investigation
 - Per-node concurrency
 - Coalescing tasks



Pros and cons: user scripts


- Users can plug in shell scripting in several places
 - In runscript, before & after postprocessing
- Enables scientific exploration
- Enables great flexibility within FRE
- Often terrible for the workflow, if best practices are not followed
- Error catching/handling is difficult




Aiming for disk affinity



- Data movement can be reduced if jobs that will reuse the same source data are launched on the same host with data in the attached local disk
- Implemented to a small degree in FRE, as a high-res option
- Must tune task parallelism vs. reduced transfers
- Difficult batch scheduling task to load balance the system



job schooling



*Art by zazer kale,
redbubble.com*



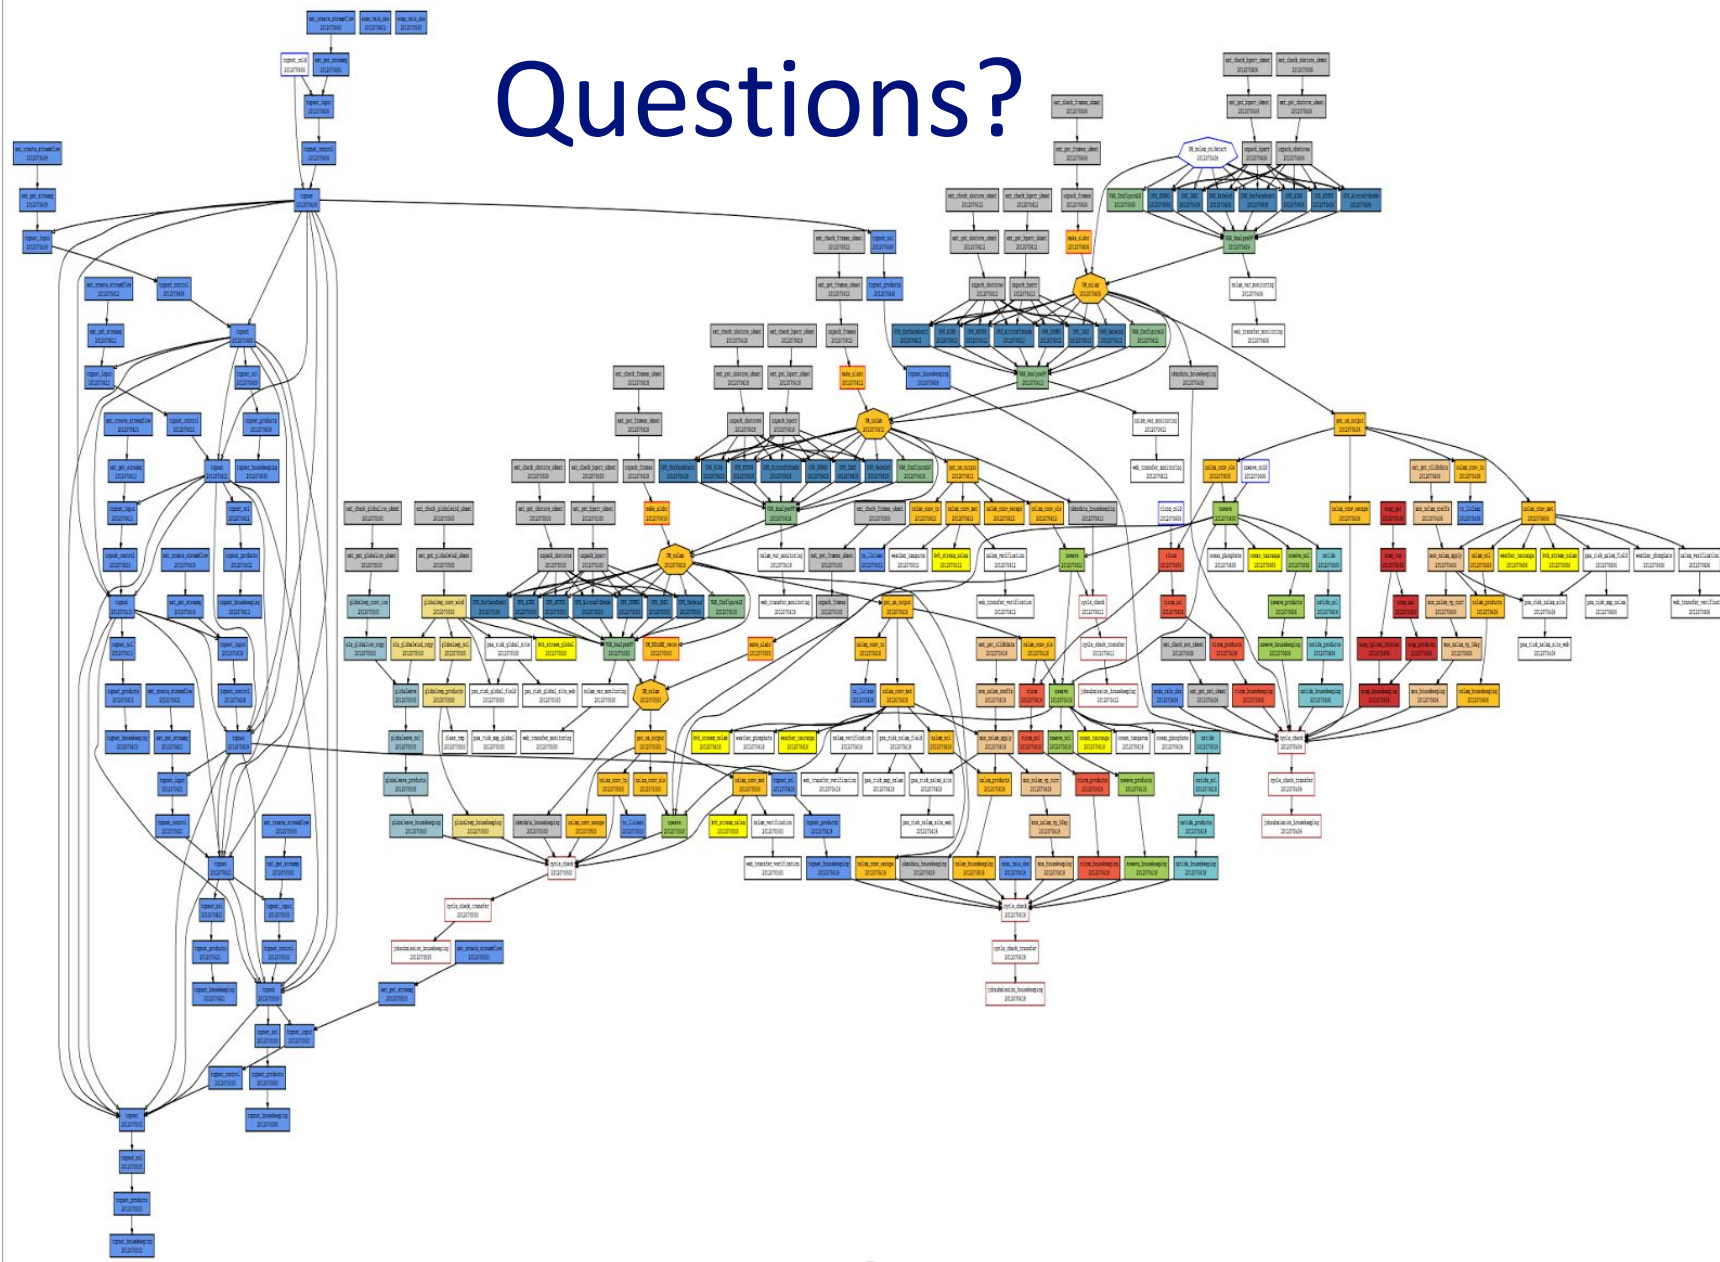


Workflow Discussion Points

- How can we reduce data movement?
- How much disk affinity can we achieve?
- Robustness: can we detect / recover from new error modes?
- How can we better leverage each other's work?
 - Automation vs. HPC security concerns
- How can we manage user plug-ins to the workflow?
- Do you trust your state files or check the filesystem?
- How do we handle user changes during a run?




Questions?





Answers?



Sample XML: Setup

```
<experimentSuite>
  <setup>
    <platform name="gfdl.default">
      <directory stem="subdirectory/paths"/>
      <csh>
        module load intel
        module load fre
      </csh>
      <property name="FMS_ARCHIVE_ROOT"
                value="/archive/fms"/>
    </platform>
  </setup>
  <experiment name=""> ...
```

Sample XML: Experiment

```
<experiment name="CM2.5A_2" inherit="CM2.5">
  <component name="fms">
    <source vc="cvs" root="/home/fms/cvs">
      <codeBase version="riga">shared</codeBase>
      <csh>
        cvs up -r riga_ar1 shared/.../file
      </csh>
    </source>
    <compile>
      <cppDefs> -Duse_netCDF </cppDefs>
    </compile>
  </component>
  <component name="mom4p1"> ...
```

Sample XML: Input

```
<input>
  <namelist name="vert_diff_driver_nml">
    do_conserve_energy = .true.
  </namelist>
  <dataFile label="input" target="INPUT/"
            chksum="" size="" timestamp="">
    $(FMS_ARCHIVE_ROOT)/am2/cover_type_field
  </dataFile>
  <diagTable>
    diagnostic variable table
  </diagTable>
</input>
```


Sample XML: Runs and Jobs

```
<runtime>
  <production simTime="26" units="years"
              npes="120" runTime="12:00:00">
    <segment simTime="12" units="months"
              runTime="06:00:00"/>
  </production>
  <regression name="basic">
    <run days="8" npes="60" runTime="00:15:00">
  </regression>
  <dataFile label="reference">
    $(FMS_ARCHIVE_ROOT)/am2/riga/19950108.tar
  </dataFile>
</runtime>
```

Sample XML: Postprocessing

```
<postProcess>
  <component type="atmos" source="atmos_month"
    sourceGrid="atmos-cubedsphere"
    xyInterp="90,144" zInterp="era40">
    <timeSeries freq="monthly"
      chunkLength="5yr">
      <variables> precip, temp </variables>
      <analysis script="/home/template.csh"/>
    </timeSeries>
    <timeAverage source="monthly"
      interval="5yr"/>
  </component>
</postProcess>
```

Archived Output

```
/archive/$USER/fre/CM2.1U_Control-1990_E1.M_3A
|-- ascii
|-- restart
|-- history
`-- pp
    |-- atmos
    |   |-- ts
    |   |   |-- monthly
    |   |   |-- 5yr
    |   |       |-- atmos.000101-000512.precip.nc
    |   |       |-- atmos.000601-001012.temp.nc
    |   |-- av
    |       |-- monthly_5yr
    |           |-- atmos.0001-0005.01.nc
    |           |-- atmos.0001-0005.02.nc
    ...
```